

CTA SCAN BLOOD CLOT DETECTION

Final Presentation



ANDROMEDA
MEDICAL IMAGING INC.

Hshmat Sahak
Nabil Mohamed
Haani Ahmed
Alaap Grandhi
Yawar Ashraf

CTA SCAN BLOOD CLOT DETECTION

Final Presentation



ANDROMEDA
MEDICAL IMAGING INC.

Hshmat Sahak
Nabil Mohamed
Haani Ahmed
Alaap Grandhi
Yawar Ashraf

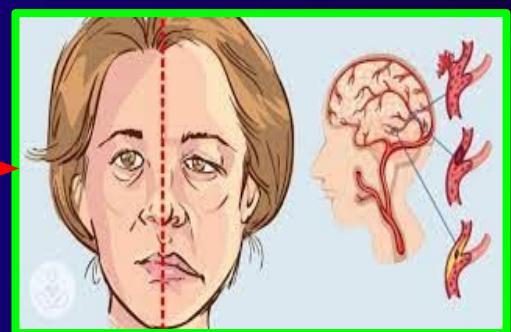
Strokes



Blood Clot blocks Blood Vessel



Parts of Brain Damaged



Disabilities like Premature Aging

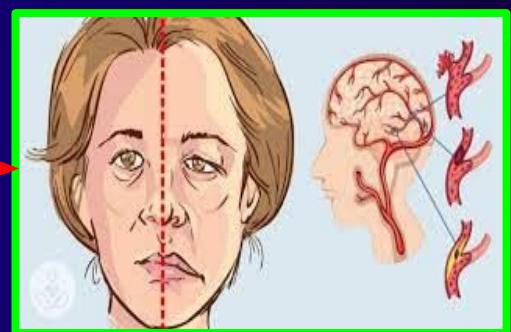
Strokes



Blood Clot blocks Blood Vessel

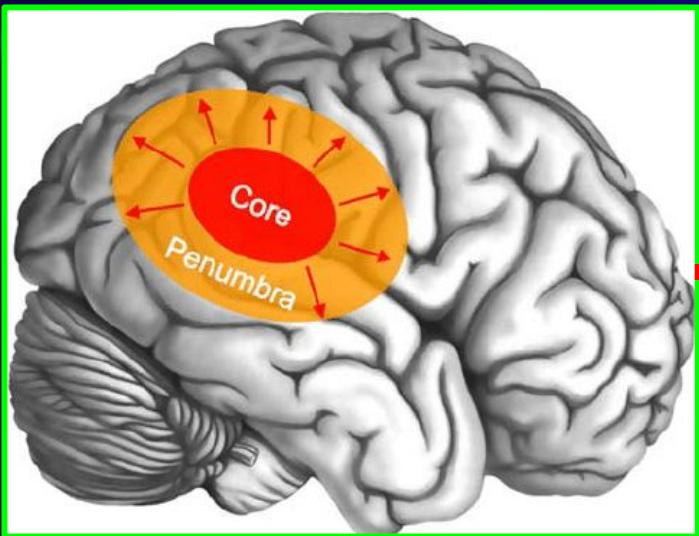


Parts of Brain Damaged

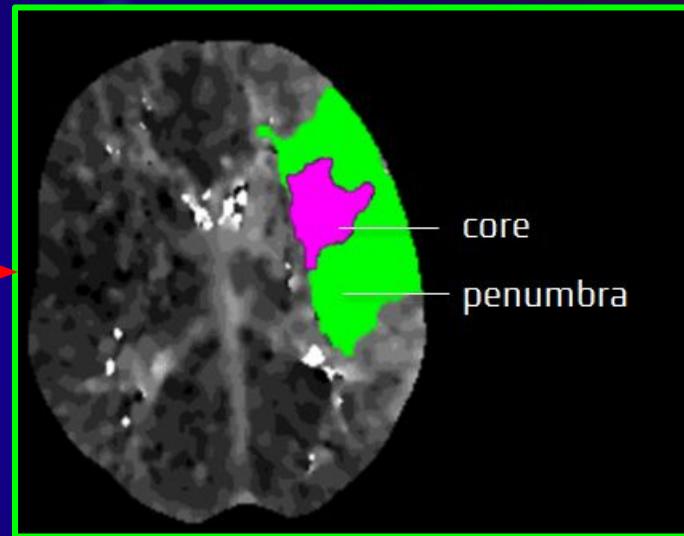


Disabilities like Premature Aging

Importance of Imaging



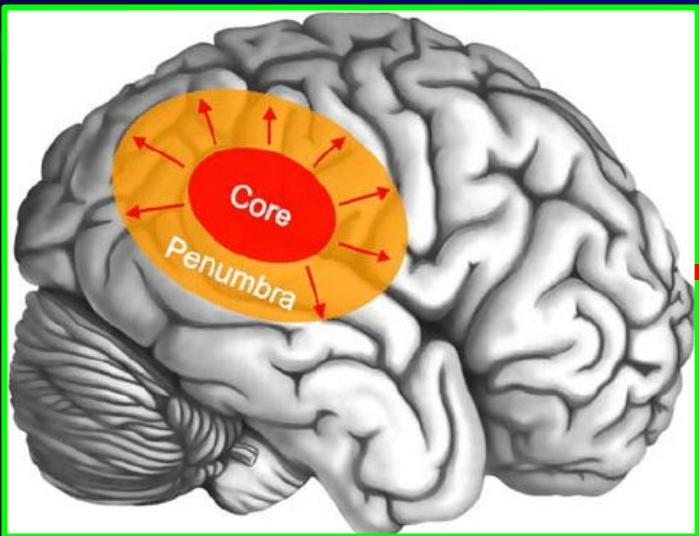
Brain with Blood Clot



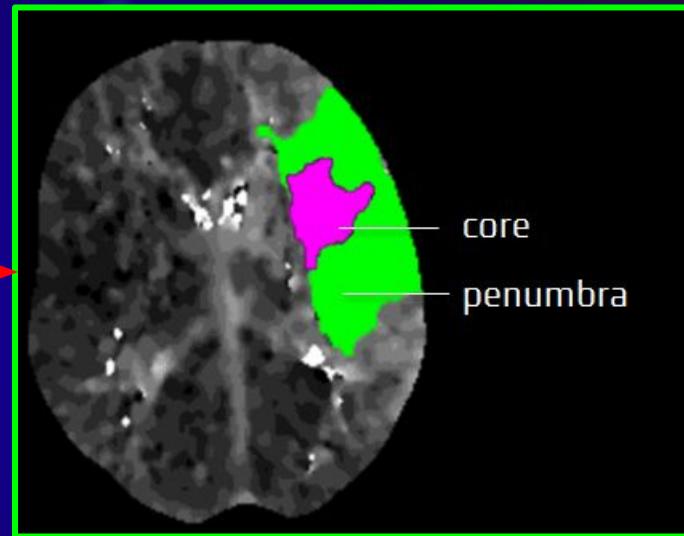
Result of Medical Imaging

- **Penumbral** = tissue affected but not yet dead
- **Core** = tissue that is already dead

Importance of Imaging



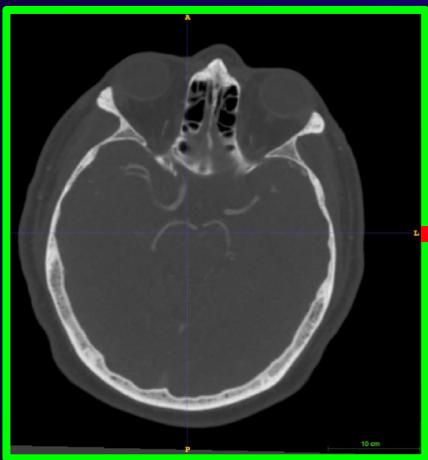
Brain with Blood Clot



Result of Medical Imaging

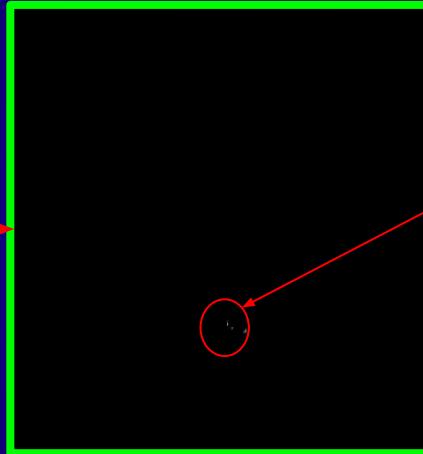
- **Penumbral** = tissue affected but not yet dead
- **Core** = tissue that is already dead

Motivation



Data from CTA Scan

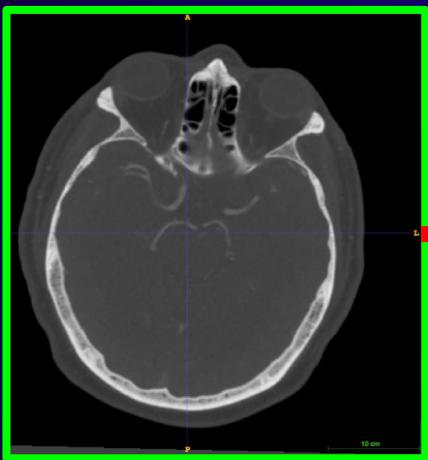
AUTOMATED method,
meant for RURAL areas



Goal: Identify Blood
Clot Region

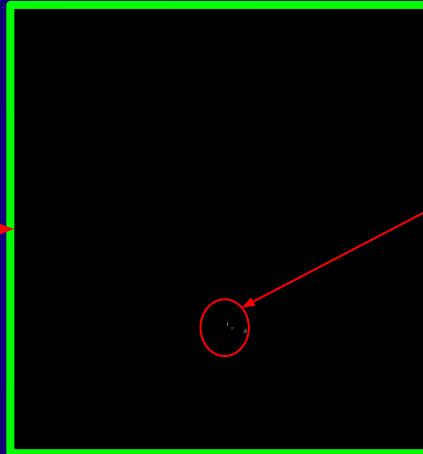
Clot

Motivation



Data from CTA Scan

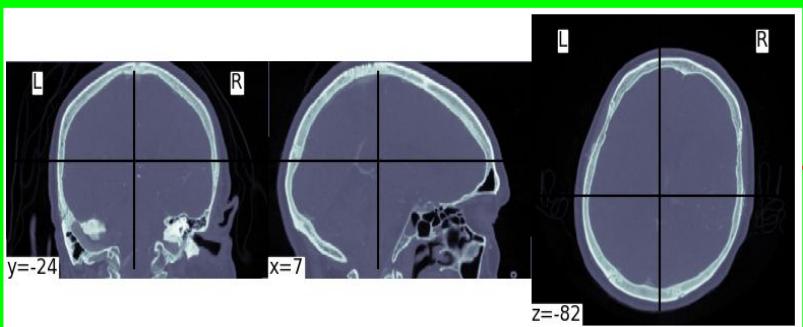
AUTOMATED method,
meant for RURAL areas



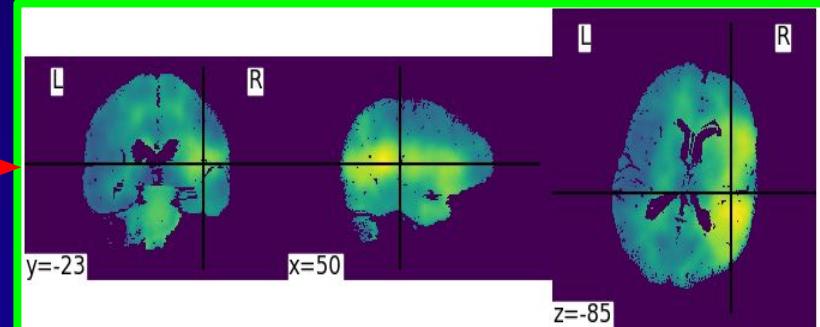
Goal: Identify Blood
Clot Region

Clot

Problem Statement

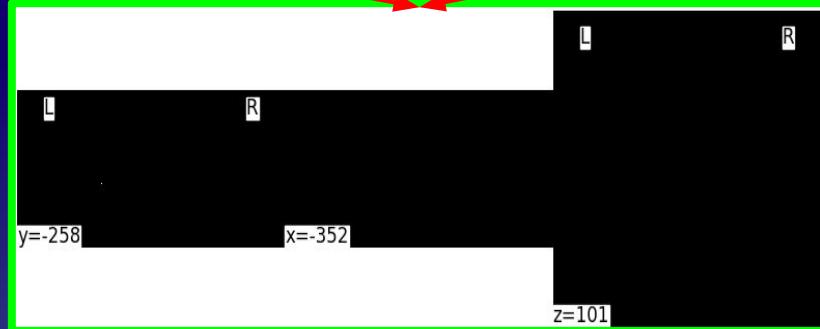


Raw mCTA Scan



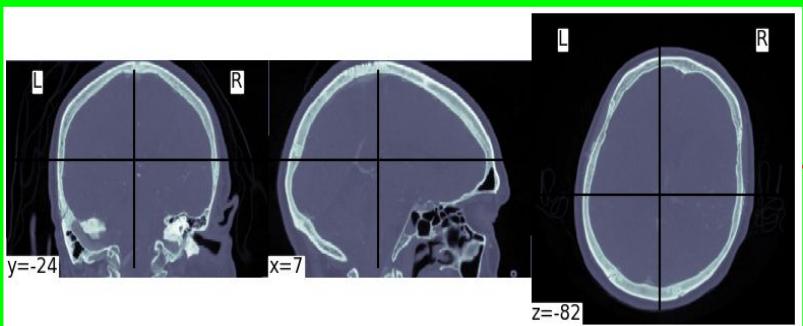
SPIRAL Output

AMI-Net

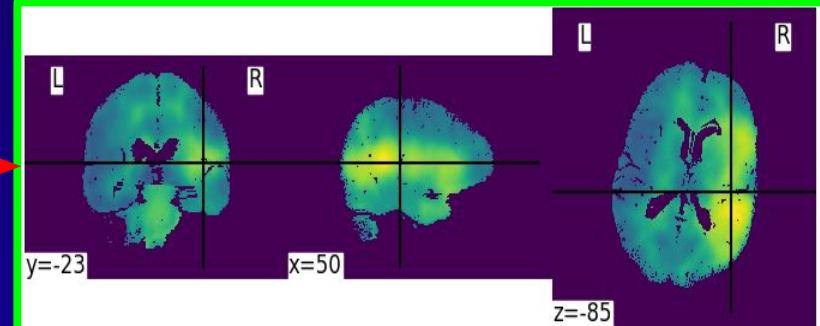


Blood Clot Region

Problem Statement

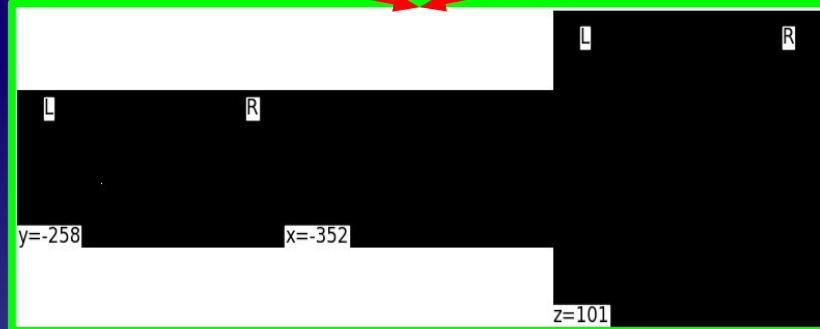


Raw mCTA Scan



SPIRAL Output

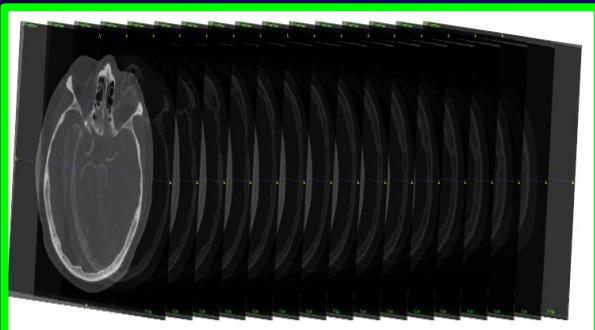
AMI-Net



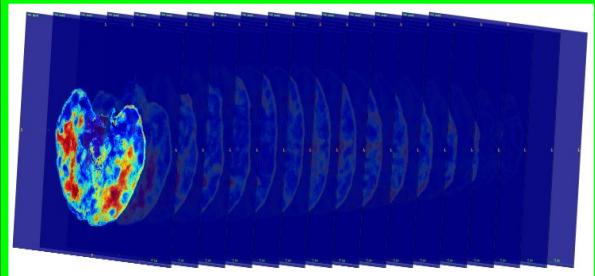
Blood Clot Region

Functional Requirements

Input Format



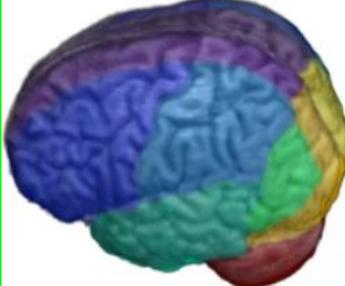
raw mCTA scans



generated SPIRAL images

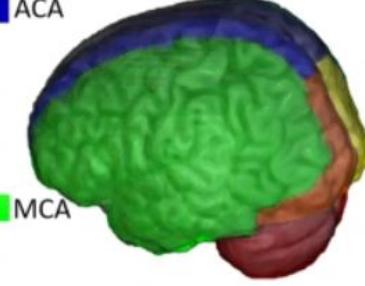
Output Format

Level 1



- ACA
- ACA medial lenticulostriate
- MCA lateral lenticulostriate
- MCA frontal
- MCA parietal
- MCA temporal
- MCA occipital
- MCA insula

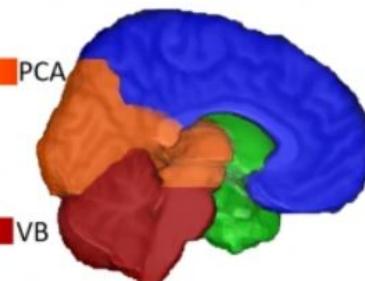
Level 2



- ACA
- MCA



- PCA temporal
- PCA occipital
- PCA posterior thalamoperforat.
- PCA anterior thalamoperforat.
- basilar
- Superior cerebellar
- Inferior cerebellar

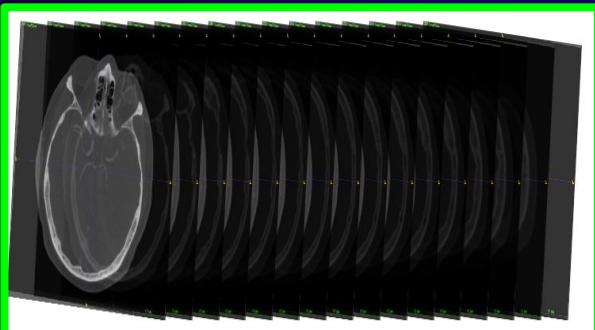


- PCA
- VB

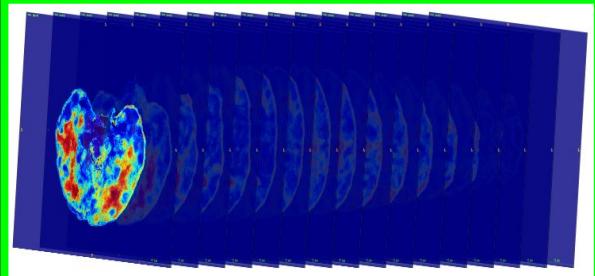
Must determine categorical site(s) of occlusion

Functional Requirements

Input Format



raw mCTA scans



generated SPIRAL images

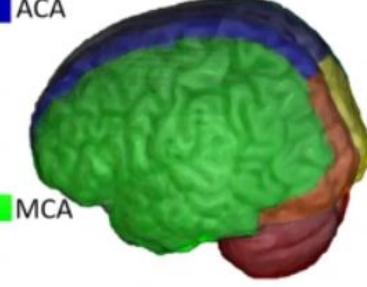
Output Format

Level 1



- ACA
- ACA medial lenticulostriate
- MCA lateral lenticulostriate
- MCA frontal
- MCA parietal
- MCA temporal
- MCA occipital
- MCA insula

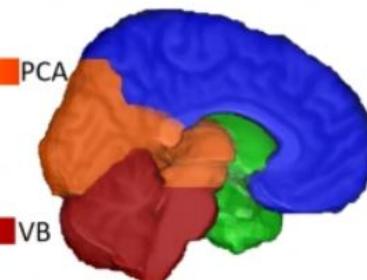
Level 2



- ACA
- MCA



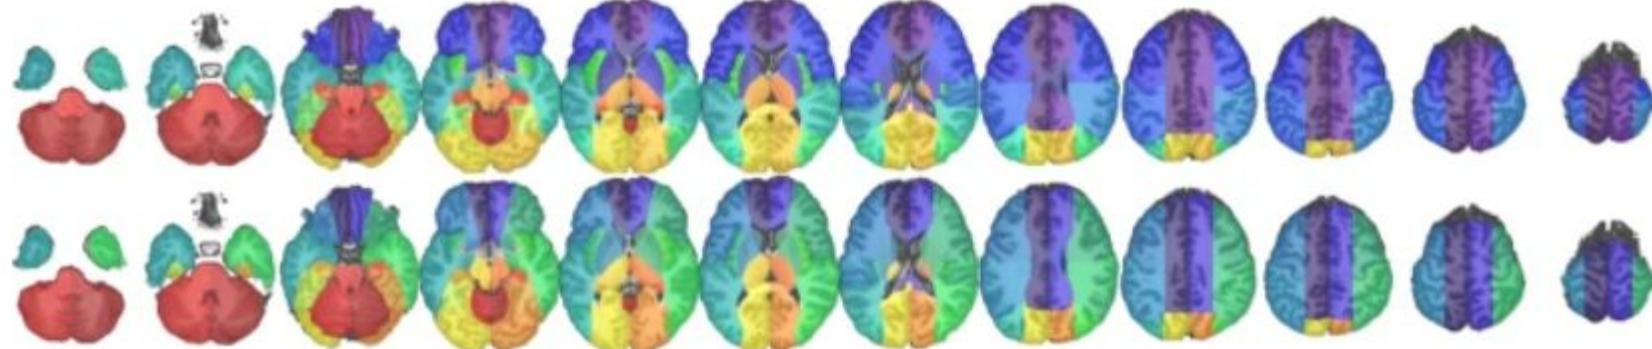
- PCA temporal
- PCA occipital
- PCA posterior thalamoperforat.
- PCA anterior thalamoperforat.
- basilar
- Superior cerebellar
- Inferior cerebellar



- PCA
- VB

Must determine categorical site(s) of occlusion

Objectives and Constraints



The blood clot will affect one or more of these 15 coloured regions

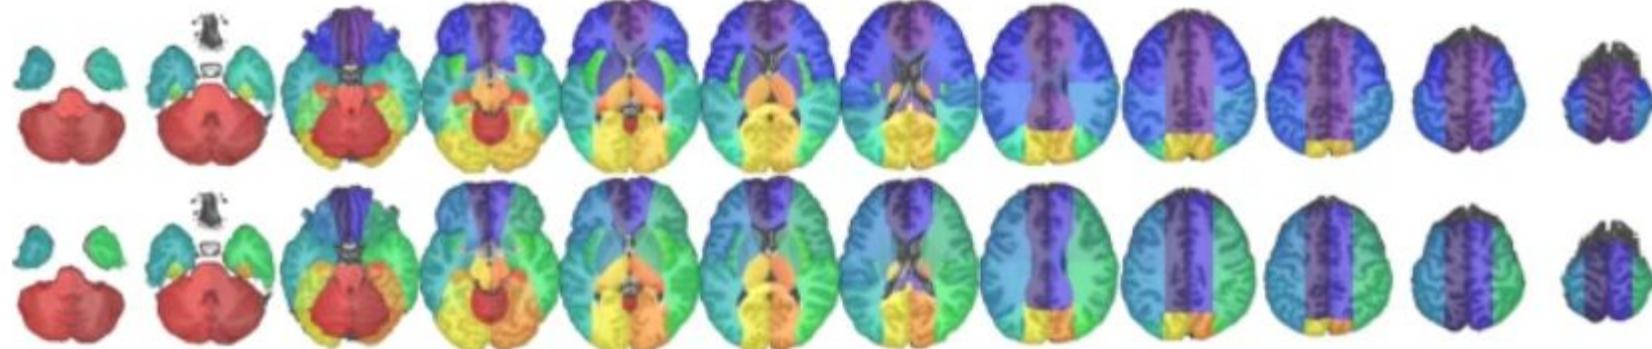
Objectives:

- *Should outperform existing SPIRAL output in identifying cranial vessels affected by clot*
- *Should identify the stroke site in a short amount of time*
- *Should have a low memory footprint*

Constraints:

- *Must run in under 5 minutes for a single patient*
- *Must be able to run on a single CPU (8-16 GB)*

Objectives and Constraints



The blood clot will affect one or more of these 15 coloured regions

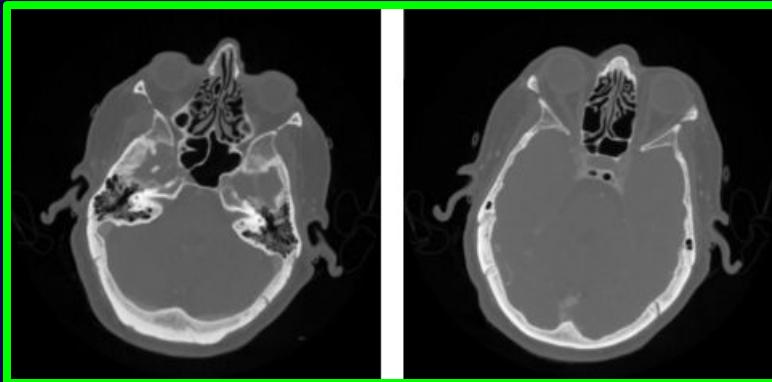
Objectives:

- *Should outperform existing SPIRAL output in identifying cranial vessels affected by clot*
- *Should identify the stroke site in a short amount of time*
- *Should have a low memory footprint*

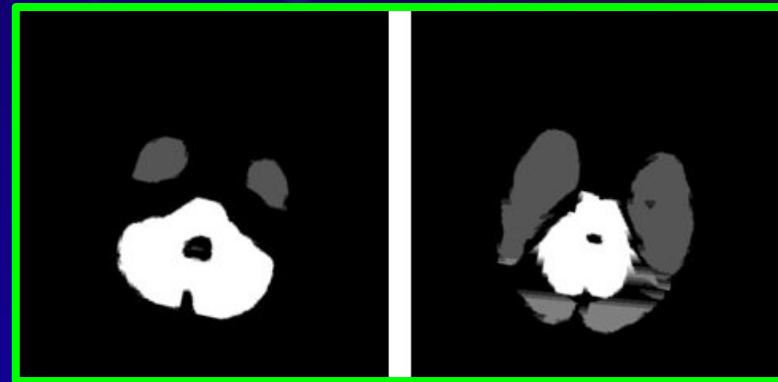
Constraints:

- *Must run in under 5 minutes for a single patient*
- *Must be able to run on a single CPU (8-16 GB)*

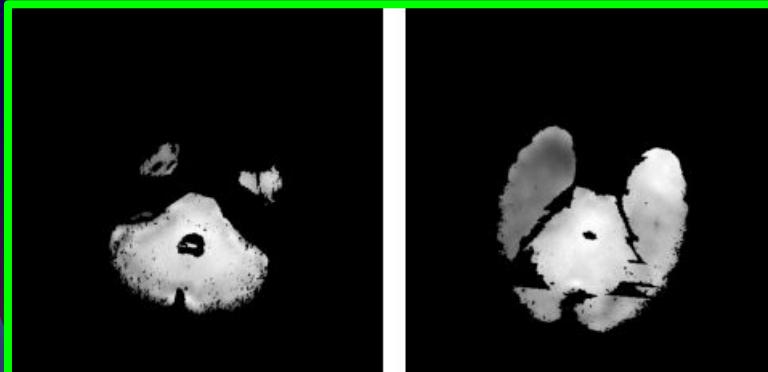
Data Visualization



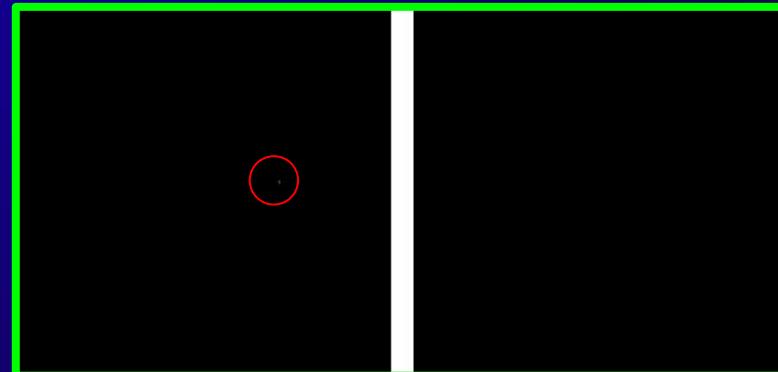
Raw CTA scan samples



ATLAS Map of Brain Regions [0-14]

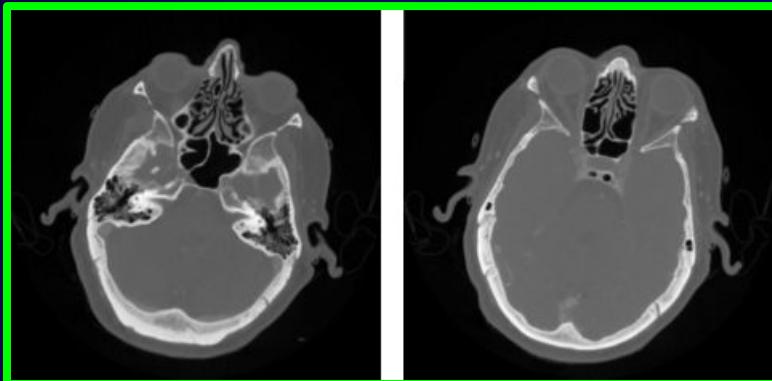


SPIRAL output - each pixel in [0,1]



Blood Clot - each pixel in [0,2]

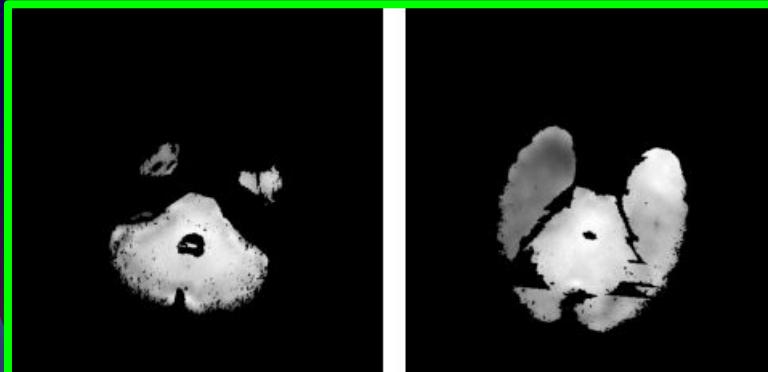
Data Visualization



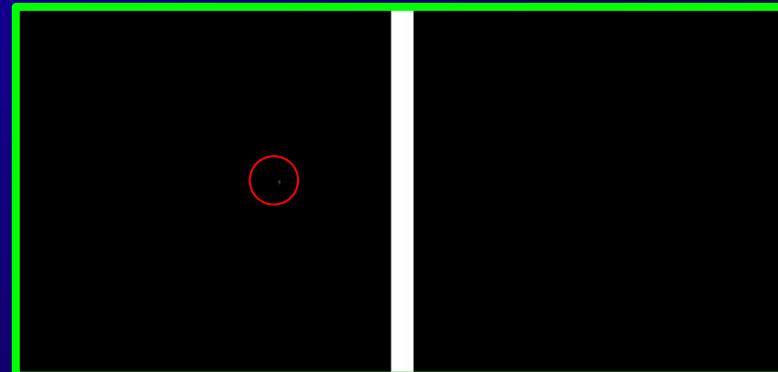
Raw CTA scan samples



ATLAS Map of Brain Regions [0-14]

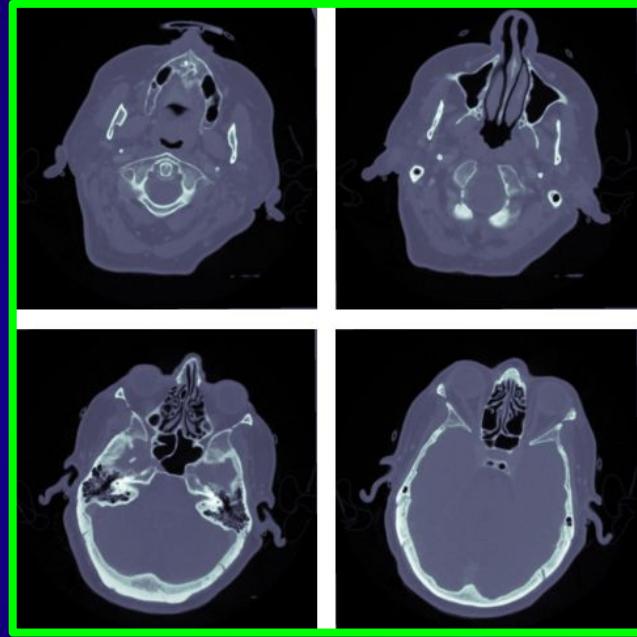
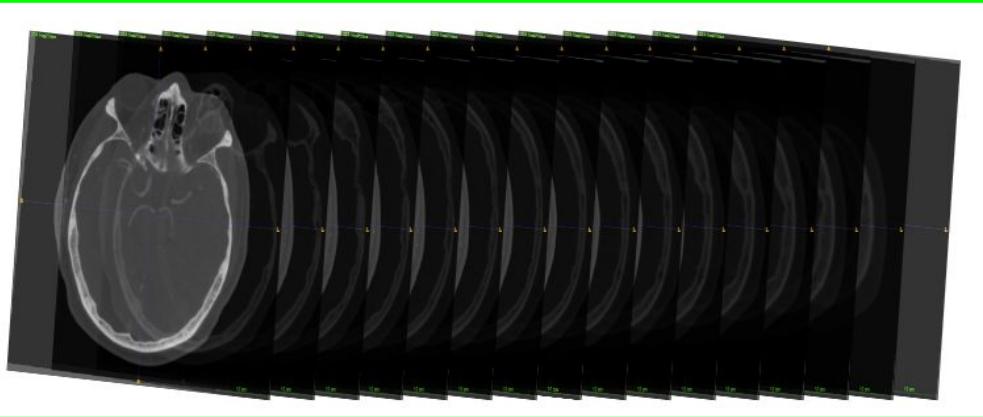


SPIRAL output - each pixel in [0,1]



Blood Clot - each pixel in [0,2]

Data Quantity and Dimensions

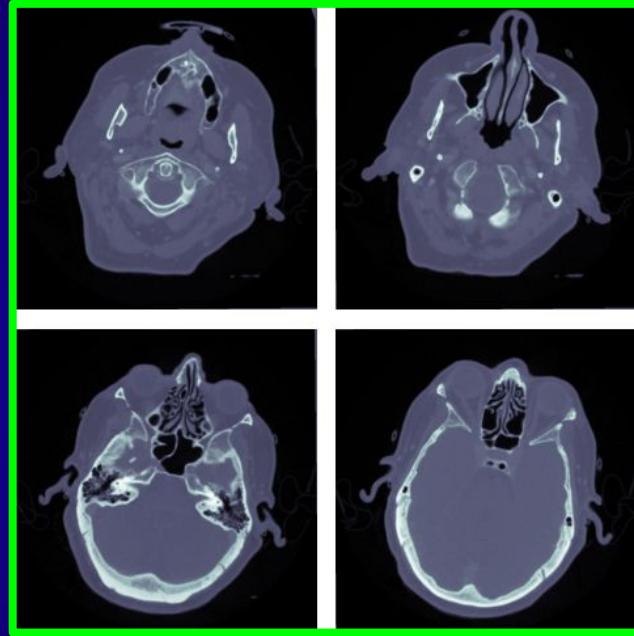
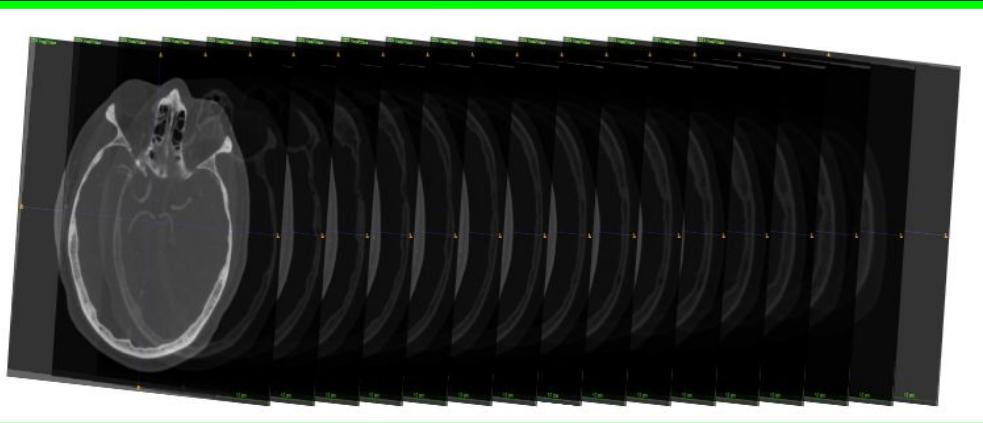


NiBabel

Access a cacophony of neuro-imaging file formats

- CTA scans are stacks of 2d images of brain at different depths
- These are in .NII (nifty) format, are 512 by 512 pixels in dimension, and vary in depth (usually ~200) by patient
- Each file is 50-70 MB

Data Quantity and Dimensions



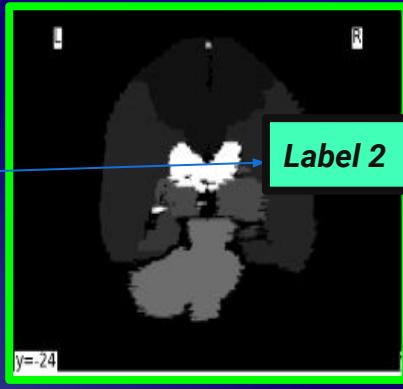
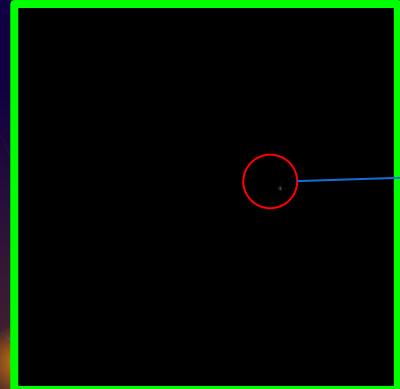
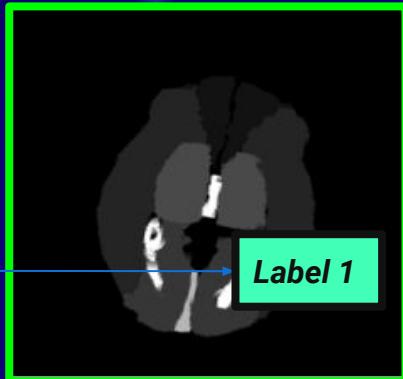
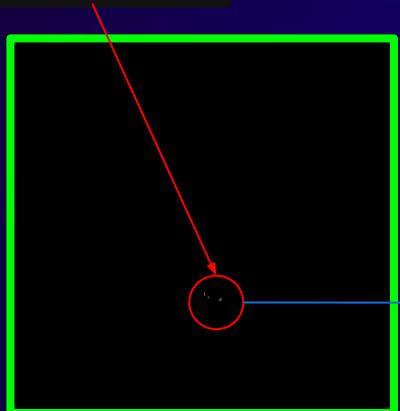
NiBabel

Access a cacophony of neuro-imaging file formats

- CTA scans are stacks of 2d images of brain at different depths
- These are in .NII (nifty) format, are 512 by 512 pixels in dimension, and vary in depth (usually ~200) by patient
- Each file is 50-70 MB

Data Labelling

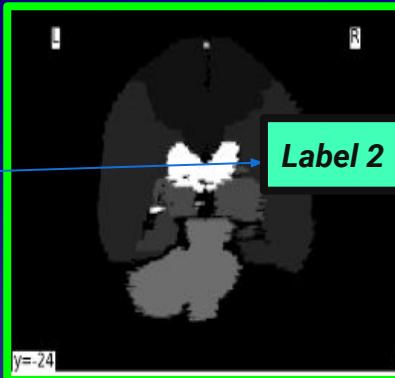
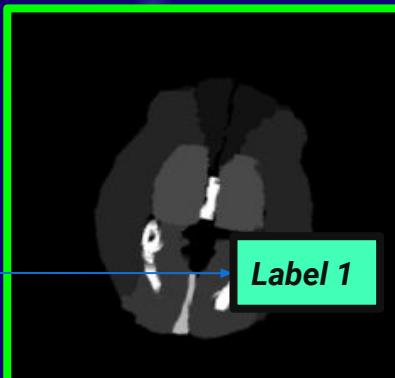
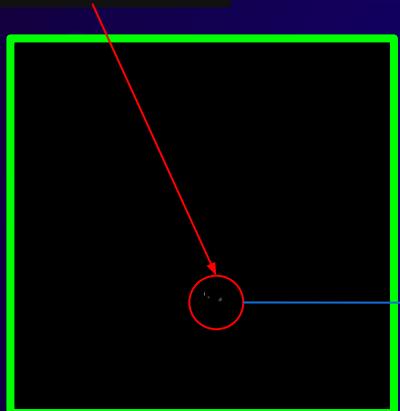
Clot in frame 1



Admire the clots

Data Labelling

Clot in frame 1



Admire the clots

Success of Project

Objective: outperform existing SPIRAL output in identifying cranial vessels affected by clot

Only operate on regions with blood clot

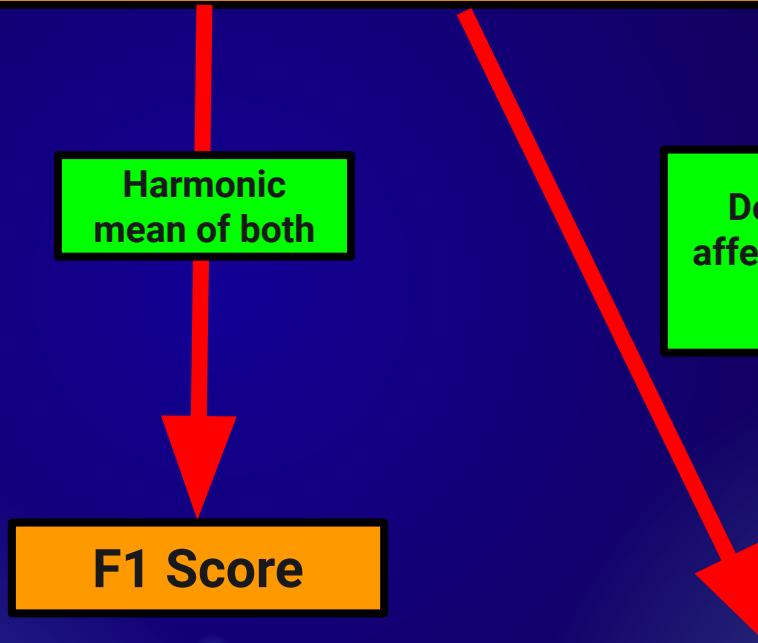
Harmonic mean of both

Don't miss any affected regions in surgery

F1 Score

Precision

Recall



Success of Project

Objective: outperform existing SPIRAL output in identifying cranial vessels affected by clot

Only operate on regions with blood clot

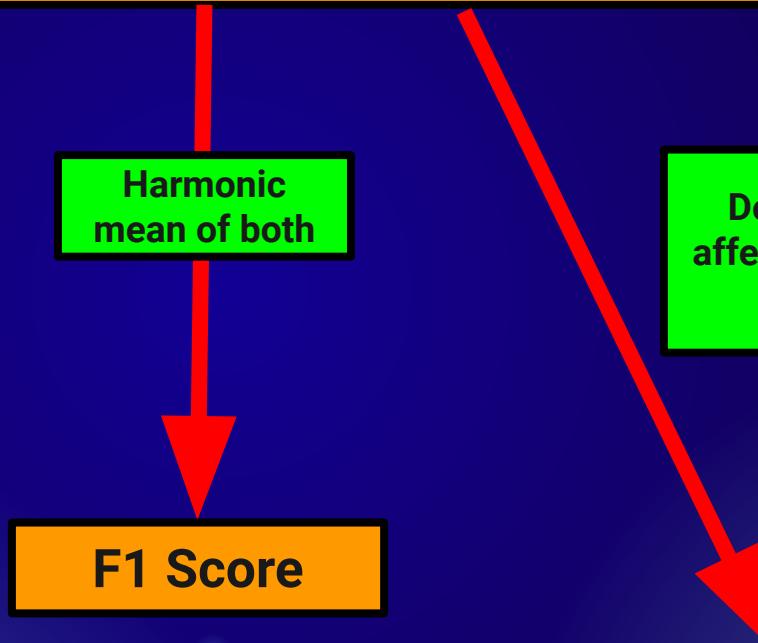
Harmonic mean of both

Don't miss any affected regions in surgery

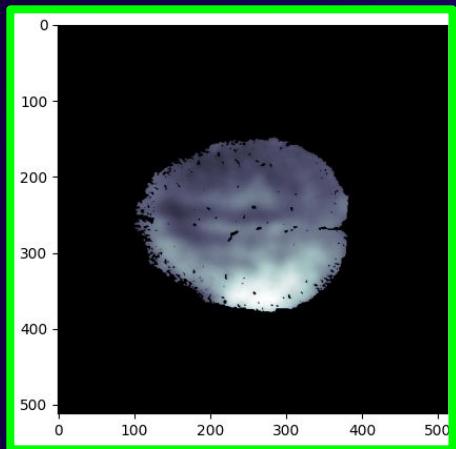
F1 Score

Precision

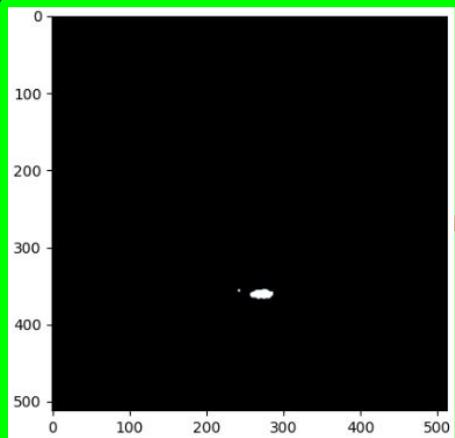
Recall



SPIRAL Data => Labels

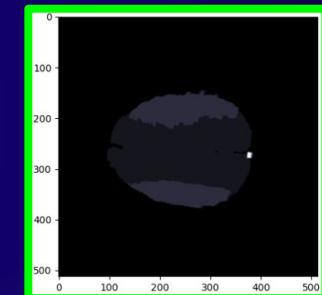


Using
Threshold



Start with SPIRAL
Frame

SPIRAL Mask (higher is
more likely penumbra)

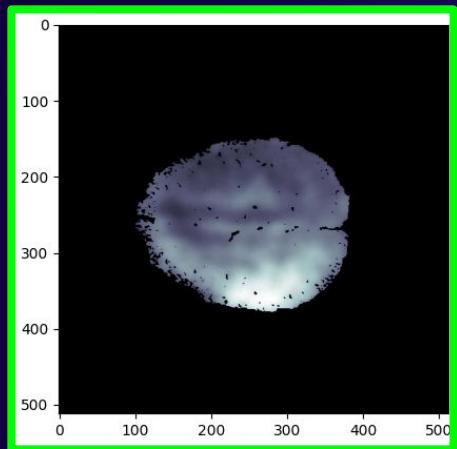


ATLAS Map

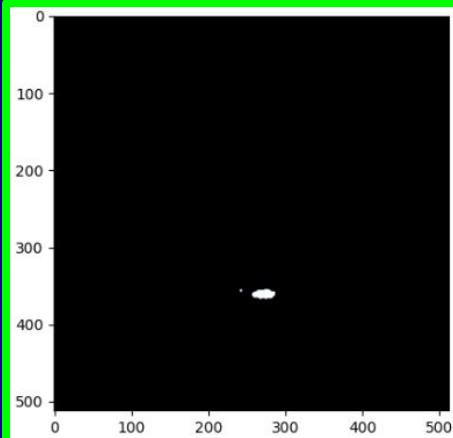
Corresponding ATLAS Regions: [0. 0. 1. 1. 1. 0.
0. 0. 0. 0. 0. 0. 0. 0. 0.]



SPIRAL Data => Labels

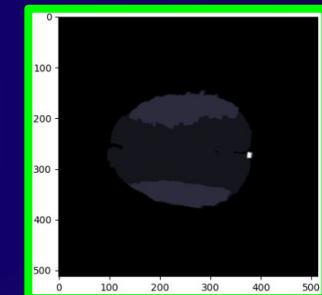


Using
Threshold



Start with SPIRAL
Frame

SPIRAL Mask (higher is
more likely penumbra)



ATLAS Map

Corresponding ATLAS Regions: [0. 0. 1. 1. 1. 0.
0. 0. 0. 0. 0. 0. 0. 0. 0.]



SPIRAL Baseline Results

```
Threshold: 0.0. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.1. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.2. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.3. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.4. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.5. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.6. Prec: 0.1558544303797345. Recall: 0.6746575342463443. Accuracy: 0.23300330033001762. F1 Score: 0.25321336757873314  
Threshold: 0.7. Prec: 0.16593886462880647. Recall: 0.6506849315066265. Accuracy: 0.30231023102308235. F1 Score: 0.2644398051171984  
Threshold: 0.8. Prec: 0.18794326241132528. Recall: 0.544520547945019. Accuracy: 0.45874587458742844. F1 Score: 0.27943760980362825  
Threshold: 0.9. Prec: 0.33439490445849224. Recall: 0.3595890410957673. Accuracy: 0.7386138613860899. F1 Score: 0.34653465341529804  
Threshold: 0.95. Prec: 0.5443037974676654. Recall: 0.1472602739725523. Accuracy: 0.8118811881187583. F1 Score: 0.23180592988549342
```

We find the optimal threshold on a train set

```
Threshold: 0.9. Prec: 0.3281356496258744. Recall: 0.3529586743659843. Accuracy: 0.7286597436512987. F1 Score: 0.3369857469852367
```

Assess the threshold found in training on test set

SPIRAL Baseline Results

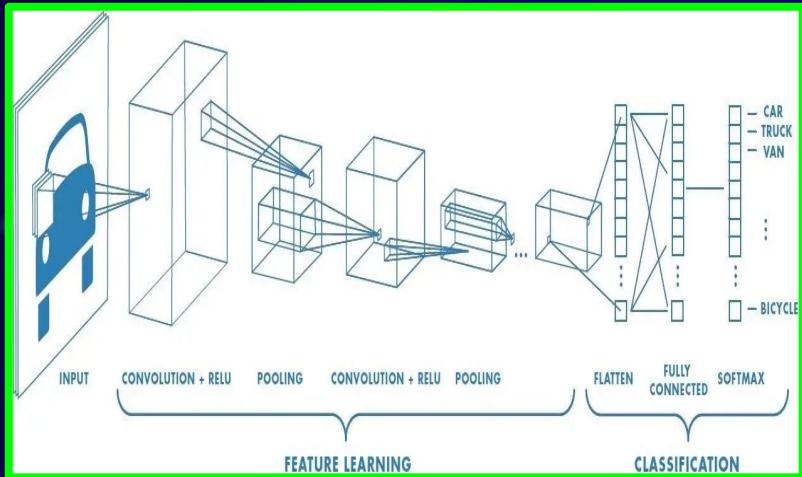
```
Threshold: 0.0. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.1. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.2. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.3. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.4. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.5. Prec: 0.15765422696114564. Recall: 0.7089041095887983. Accuracy: 0.21386138613859973. F1 Score: 0.25794392520384624  
Threshold: 0.6. Prec: 0.1558544303797345. Recall: 0.6746575342463443. Accuracy: 0.23300330033001762. F1 Score: 0.25321336757873314  
Threshold: 0.7. Prec: 0.16593886462880647. Recall: 0.6506849315066265. Accuracy: 0.30231023102308235. F1 Score: 0.2644398051171984  
Threshold: 0.8. Prec: 0.18794326241132528. Recall: 0.544520547945019. Accuracy: 0.45874587458742844. F1 Score: 0.27943760980362825  
Threshold: 0.9. Prec: 0.33439490445849224. Recall: 0.3595890410957673. Accuracy: 0.7386138613860899. F1 Score: 0.34653465341529804  
Threshold: 0.95. Prec: 0.5443037974676654. Recall: 0.1472602739725523. Accuracy: 0.8118811881187583. F1 Score: 0.23180592988549342
```

We find the optimal threshold on a train set

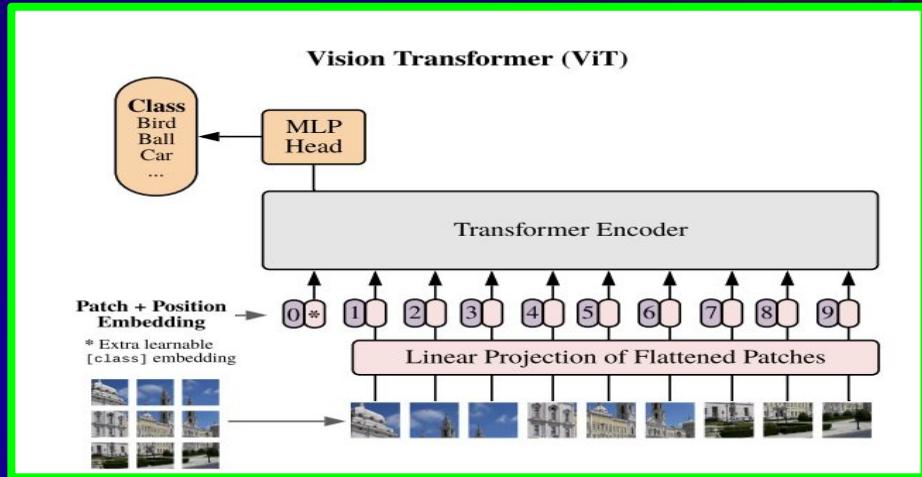
```
Threshold: 0.9. Prec: 0.3281356496258744. Recall: 0.3529586743659843. Accuracy: 0.7286597436512987. F1 Score: 0.3369857469852367
```

Assess the threshold found in training on test set

DL Computer Vision Models



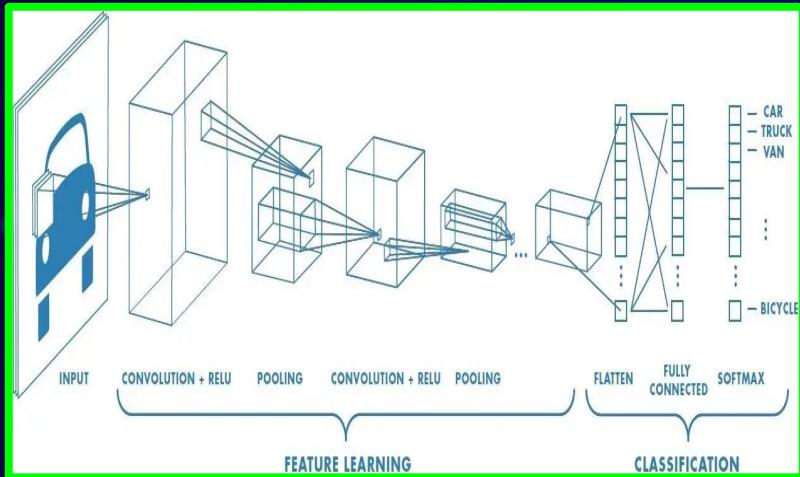
CNN



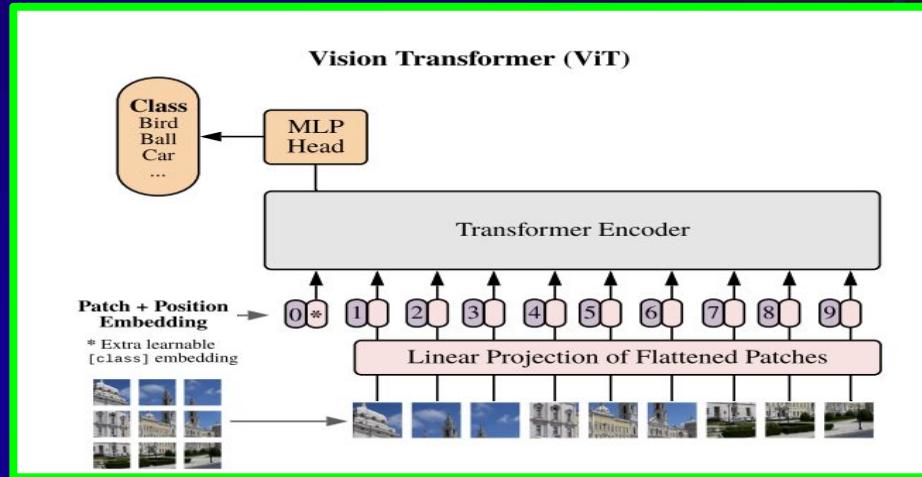
Vision Transformer

- **CNNs and Vision Transformers: State-of-the-art in various tasks in medical imaging**
- **CNNs are efficient due to fewer (shared) parameters.**
- **Vision Transformers have achieved high performance in detecting retinal diseases and analyzing images for cancer diagnosis (similar domains) but require more data.**

DL Computer Vision Models



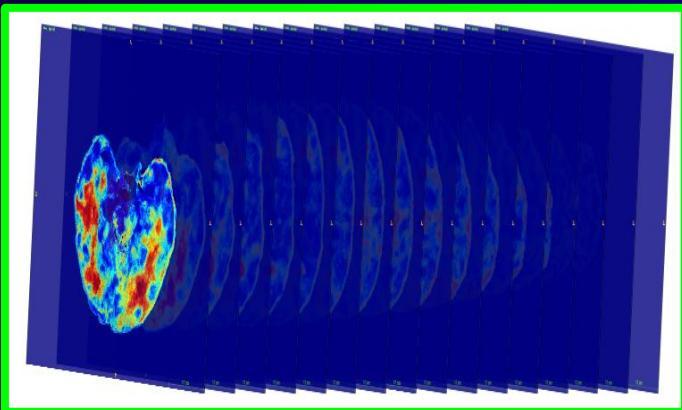
CNN



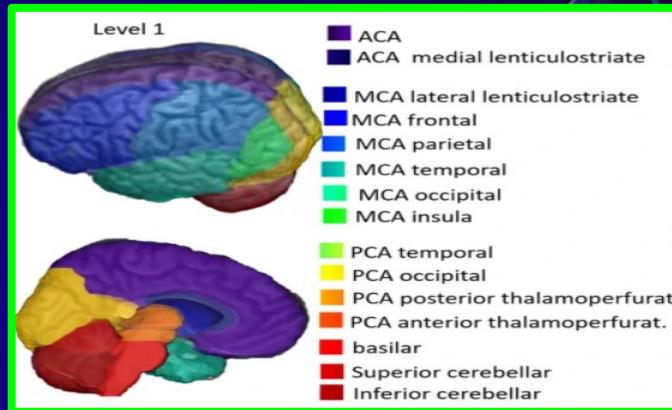
Vision Transformer

- **CNNs and Vision Transformers: State-of-the-art in various tasks in medical imaging**
- **CNNs are efficient due to fewer (shared) parameters.**
- **Vision Transformers have achieved high performance in detecting retinal diseases and analyzing images for cancer diagnosis (similar domains) but require more data.**

How do we apply DL here?



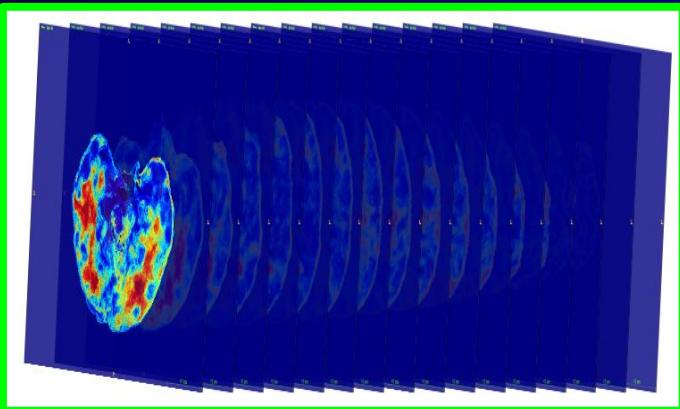
- 1 Baseline CNN
- 2 MobileNet
- 3 EfficientNet



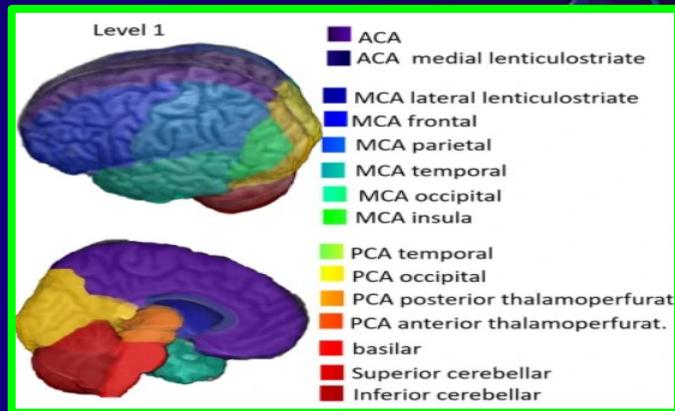
Input can be i) frame, ii) all frames
iii) sliding window of frames (e.g 5 frames per input)

Output is length 15 array of probabilities of each region being affected

How do we apply DL here?



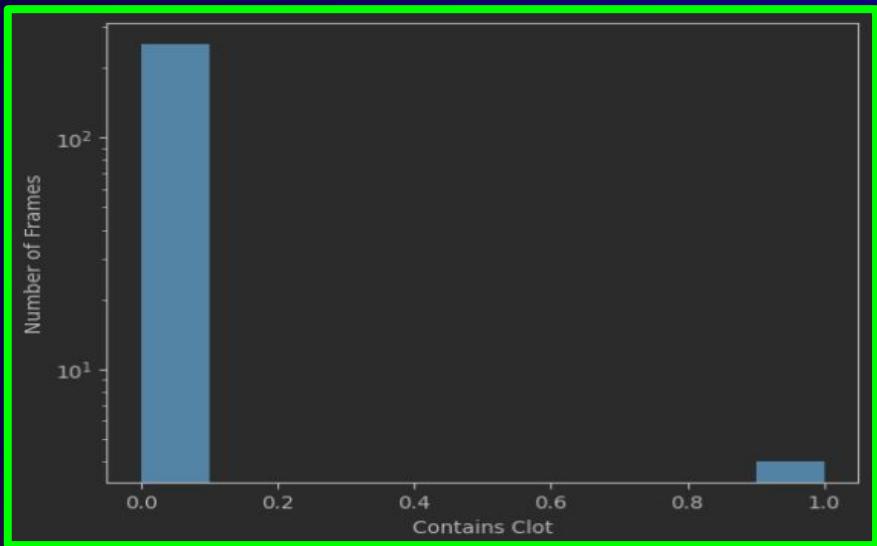
- 1 Baseline CNN
- 2 MobileNet
- 3 EfficientNet



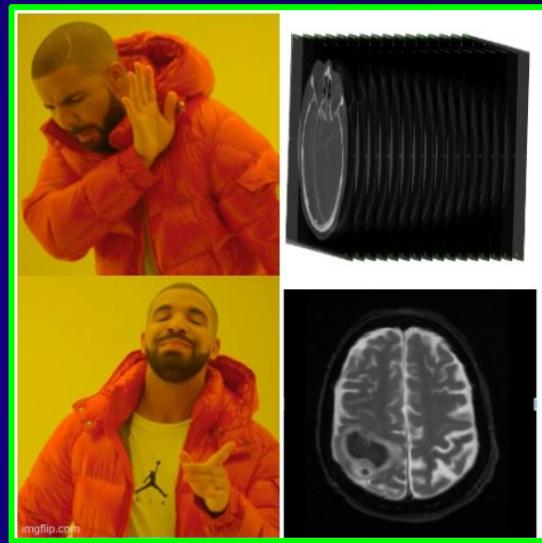
Input can be i) frame, ii) all frames
iii) sliding window of frames (e.g 5 frames per input)

Output is length 15 array of probabilities of each region being affected

Potential Roadblocks

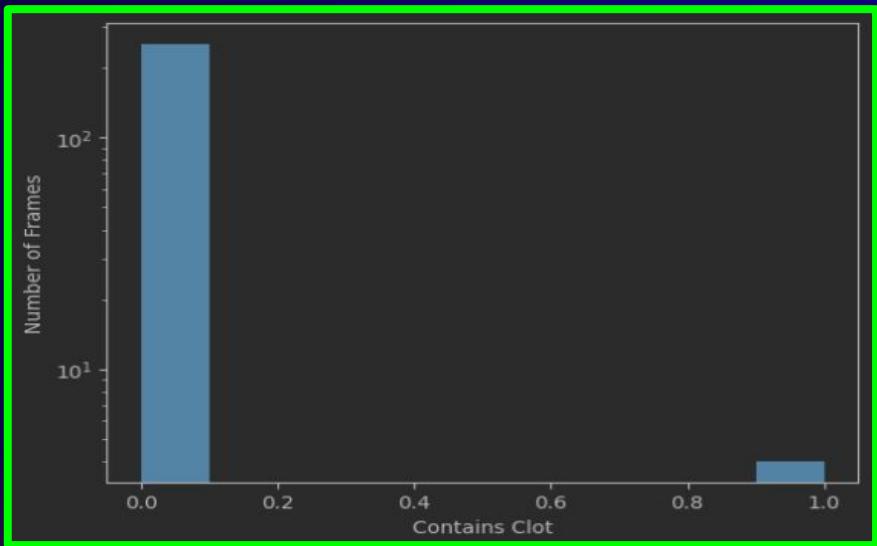


Imbalanced Dataset

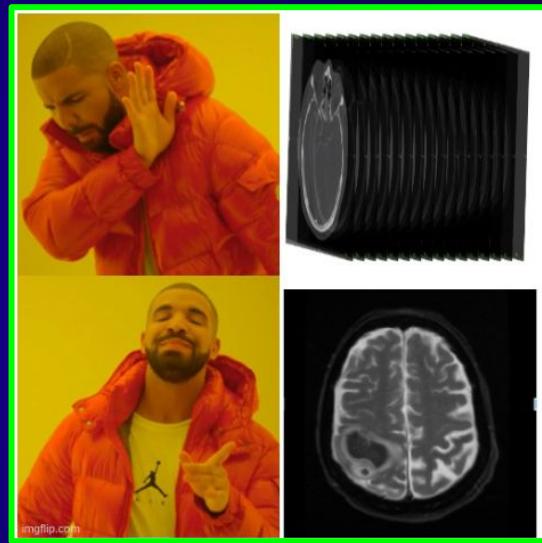


CUDA Memory Issues

Potential Roadblocks

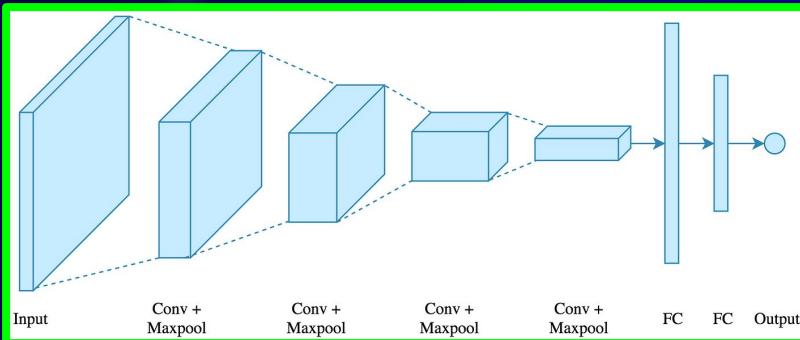


Imbalanced Dataset



CUDA Memory Issues

Model 1: Baseline CNN

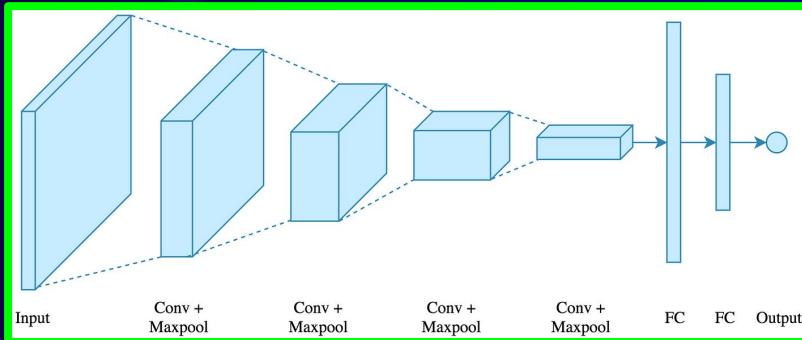


Model Architecture

Layer	No. of Kernels	Kernel Size	Activation
Input	1	Input Shape	-
Convolution2D	32	5×5	ReLU
Maxpooling2D	-	2×2	-
Convolution2D	32	5×5	ReLU
Maxpooling2D	-	2×2	-
Fully Connected	1024	-	Dropout (0.25)
Fully Connected	2	-	ReLU
SoftMax	-	-	-

The layers composing our model, taken from the paper
(Zivkovic et al., *Hybrid CNN and XGBoost Model Tuned by Modified Arithmetic Optimization Algorithm for COVID-19 Early Diagnostics from X-ray Images*, 2022)

Model 1: Baseline CNN

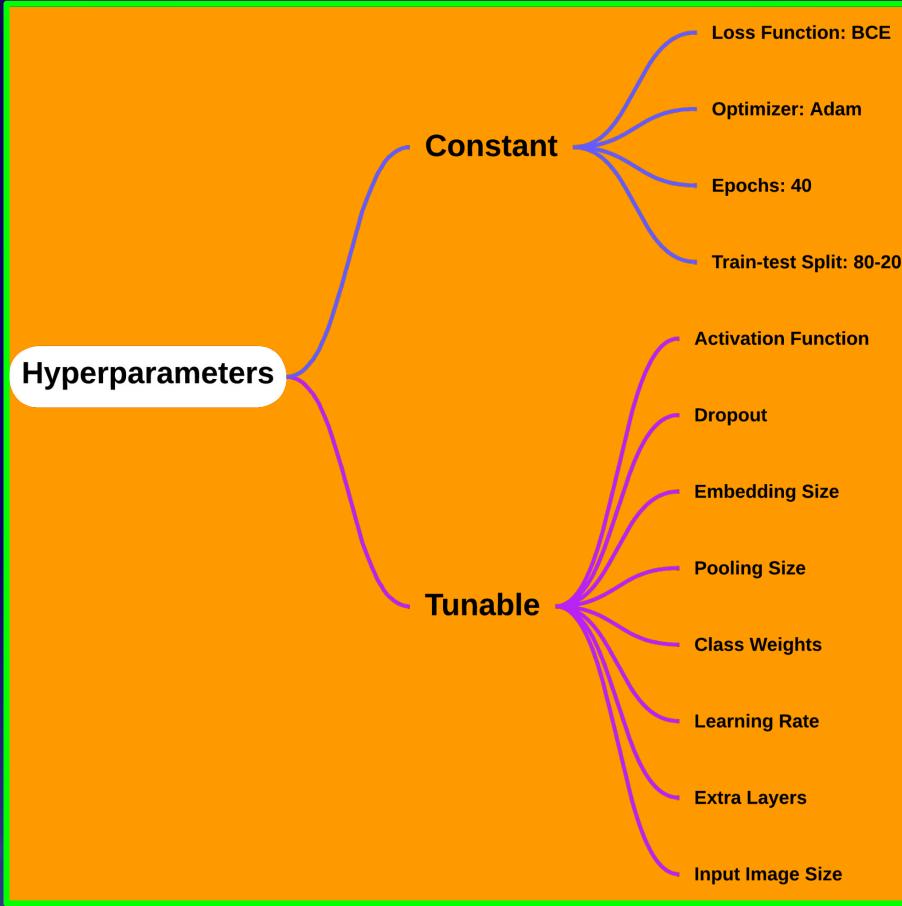


Model Architecture

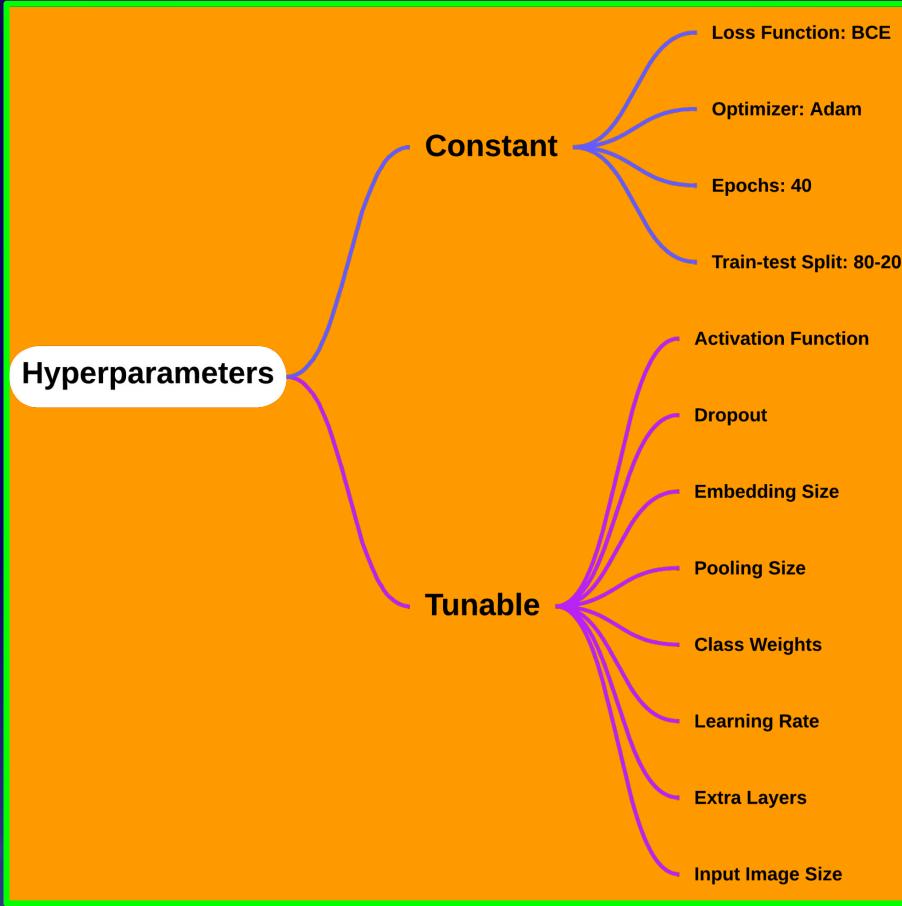
Layer	No. of Kernels	Kernel Size	Activation
Input	1	Input Shape	-
Convolution2D	32	5×5	ReLU
Maxpooling2D	-	2×2	-
Convolution2D	32	5×5	ReLU
Maxpooling2D	-	2×2	-
Fully Connected	1024	-	Dropout (0.25)
Fully Connected	2	-	ReLU
SoftMax	-	-	-

The layers composing our model, taken from the paper
(Zivkovic et al., *Hybrid CNN and XGBoost Model Tuned by Modified Arithmetic Optimization Algorithm for COVID-19 Early Diagnostics from X-ray Images*, 2022)

Model Optimization



Model Optimization



Hyperparameter Tuning

pool_dim	feat_dim	activation	lr	dropout_prob	pos_weight
4	256	relu	0.000512239	0.340438	17.428
4	256	sigmoid	0.00107452	0.181308	12.9439
4	256	tanh	0.00487844	0.307489	15.844
4	512	tanh	0.00142236	0.46978	13.3317
4	256	sigmoid	0.00234546	0.155405	13.8587
4	512	sigmoid	0.00094525	0.288003	14.2234
4	512	sigmoid	0.00067001	0.121338	15.4934
4	512	relu	0.00189603	0.123148	16.2624
4	512	relu	0.00350341	0.186474	15.2463
4	256	sigmoid	0.00210993	0.402286	16.6853

```
1 hyperparameters:  
2     activation: sigmoid  
3     batch: 128  
4     feat_dim: 32  
5     lr: 0.0001  
6     model: CNN_Feature_Extractor  
7     num_channels: 210  
8     optimizer_name: adam  
9     pooling_size: 2  
10    drop_out_rate: 0.2
```

RayTune Library for Tuning



Tuned Hyperparameters

Hyperparameter Tuning

pool_dim	feat_dim	activation	lr	dropout_prob	pos_weight
4	256	relu	0.000512239	0.340438	17.428
4	256	sigmoid	0.00107452	0.181308	12.9439
4	256	tanh	0.00487844	0.307489	15.844
4	512	tanh	0.00142236	0.46978	13.3317
4	256	sigmoid	0.00234546	0.155405	13.8587
4	512	sigmoid	0.00094525	0.288003	14.2234
4	512	sigmoid	0.00067001	0.121338	15.4934
4	512	relu	0.00189603	0.123148	16.2624
4	512	relu	0.00350341	0.186474	15.2463
4	256	sigmoid	0.00210993	0.402286	16.6853

```
1 hyperparameters:  
2     activation: sigmoid  
3     batch: 128  
4     feat_dim: 32  
5     lr: 0.0001  
6     model: CNN_Feature_Extractor  
7     num_channels: 210  
8     optimizer_name: adam  
9     pooling_size: 2  
10    drop_out_rate: 0.2
```

RayTune Library for Tuning



Tuned Hyperparameters

Hyperparameter Tuning Results

pool_dim	feat_dim	activation	lr	dropout_prob	pos_weight	extra_layers	resize_dims	iter	total time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
4	512	tanh	0.00132562	0.318361	17.7171	0	(128, 128)	8	79.8011	0.327614	0.199568	0.914191	0.963328
4	512	relu	0.000608253	0.387794	15.2612	2	(128, 128)	20	106.681	0.598461	0.459725	0.857143	0.989873
4	512	sigmoid	0.00312816	0.196748	15.5842	2	(512, 512)	4	288.843	0.184887	0.353612	0.125162	0.986776
4	512	relu	0.00480434	0.131557	12.2852	4	(512, 512)	8	478.249	0.401463	0.271881	0.767075	0.978939
4	512	sigmoid	0.000681346	0.166743	13.9148	2	(128, 128)	4	67.417	0.168809	0.106594	0.405498	0.962522

Results using SPIRAL as Input

pool_dim	feat_dim	activation	lr	dropout_prob	pos_weight	extra_layers	resize_dims	iter	total time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
4	512	tanh	0.00109909	0.248508	12.9139	2	(256, 256)	20	325.903	0.620511	0.514077	0.78254	0.990276
4	512	relu	0.00186225	0.308561	13.4125	2	(256, 256)	16	298.293	0.474085	0.322626	0.893617	0.98347
4	512	sigmoid	0.000671123	0.294085	14.9742	2	(512, 512)	4	252.012	0.191932	0.159434	0.241084	0.977052
4	512	tanh	0.000878918	0.274336	15.6468	2	(128, 128)	20	132.717	0.666351	0.514059	0.946879	0.988486
4	512	sigmoid	0.00141349	0.321913	14.2752	2	(128, 128)	4	52.2891	0.355104	0.229973	0.778963	0.970069

Results using raw CTA as Input

Hyperparameter Tuning Results

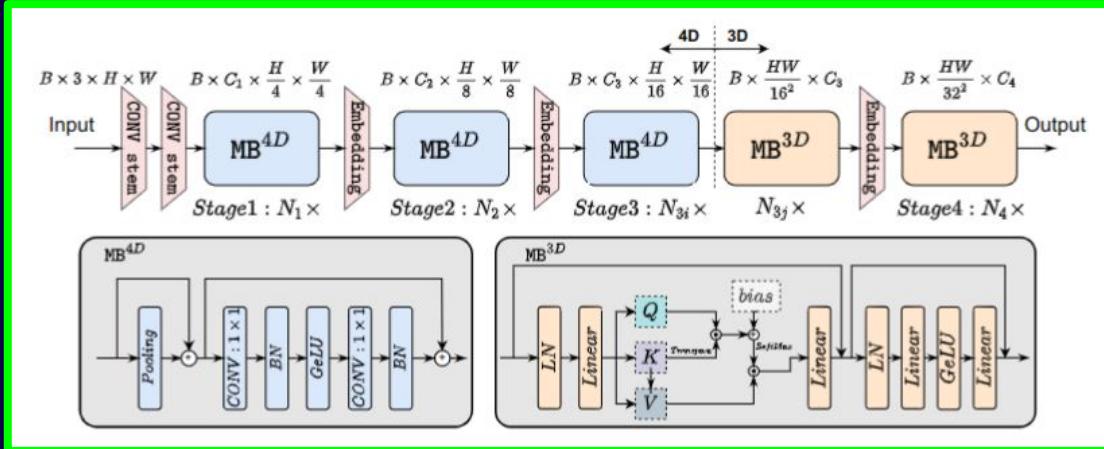
pool_dim	feat_dim	activation	lr	dropout_prob	pos_weight	extra_layers	resize_dims	iter	total time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
4	512	tanh	0.00132562	0.318361	17.7171	0	(128, 128)	8	79.8011	0.327614	0.199568	0.914191	0.963328
4	512	relu	0.000608253	0.387794	15.2612	2	(128, 128)	20	106.681	0.598461	0.459725	0.857143	0.989873
4	512	sigmoid	0.00312816	0.196748	15.5842	2	(512, 512)	4	288.843	0.184887	0.353612	0.125162	0.986776
4	512	relu	0.00480434	0.131557	12.2852	4	(512, 512)	8	478.249	0.401463	0.271881	0.767075	0.978939
4	512	sigmoid	0.000681346	0.166743	13.9148	2	(128, 128)	4	67.417	0.168809	0.106594	0.405498	0.962522

Results using SPIRAL as Input

pool_dim	feat_dim	activation	lr	dropout_prob	pos_weight	extra_layers	resize_dims	iter	total time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
4	512	tanh	0.00109909	0.248508	12.9139	2	(256, 256)	20	325.903	0.620511	0.514077	0.78254	0.990276
4	512	relu	0.00186225	0.308561	13.4125	2	(256, 256)	16	298.293	0.474085	0.322626	0.893617	0.98347
4	512	sigmoid	0.000671123	0.294085	14.9742	2	(512, 512)	4	252.012	0.191932	0.159434	0.241084	0.977052
4	512	tanh	0.000878918	0.274336	15.6468	2	(128, 128)	20	132.717	0.666351	0.514059	0.946879	0.988486
4	512	sigmoid	0.00141349	0.321913	14.2752	2	(128, 128)	4	52.2891	0.355104	0.229973	0.778963	0.970069

Results using raw CTA as Input

Model 2: EfficientFormerV2



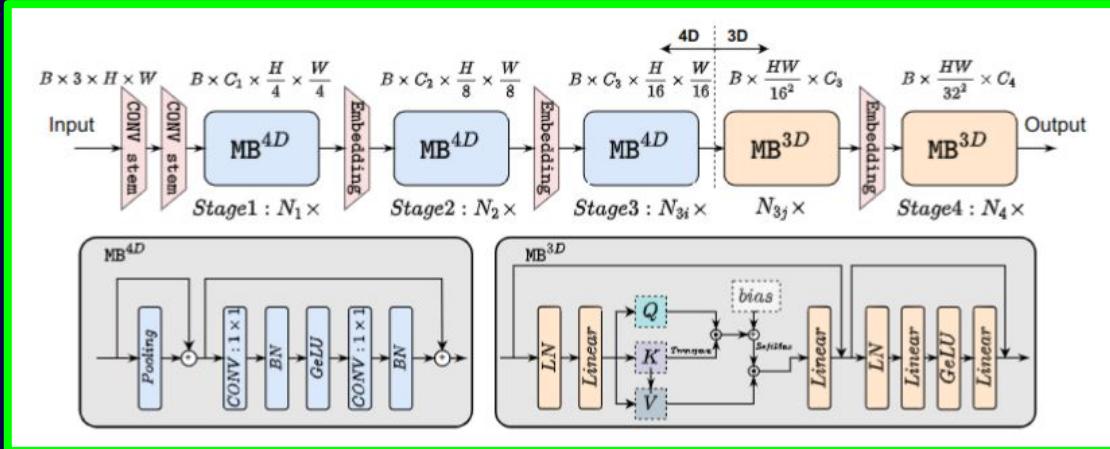
Model Architecture

```
1 hyperparameters:  
2 batch_size: 32  
3 lr: 0.0015634  
4 pos_weight: 15.9876  
5
```

Tuned Hyperparameters

- Similar to previous slide, use pre-trained EfficientFormerV2 as feature extractor
- Add classification layer on top
- Tune classifier from scratch and finetune architecture

Model 2: EfficientFormerV2



Model Architecture

```
1 hyperparameters:  
2   batch_size: 32  
3   lr: 0.0015634  
4   pos_weight: 15.9876  
5
```

Tuned Hyperparameters

- Similar to previous slide, use pre-trained EfficientFormerV2 as feature extractor
- Add classification layer on top
- Tune classifier from scratch and finetune architecture

Training Results

lr	pos_weight	resize_dims	iter	total_time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
0.000855178	16.6069 (224, 224)		1	79.4564	0.0598473	0.154839	0.0370943	0.987841
0.00132556	13.5432 (224, 224)		20	862.509	0.600391	0.438462	0.951983	0.990211
0.00199256	12.2775 (224, 224)		20	859.569	0.57637536	0.42092352	0.91390368	0.95060256
0.00109437	16.3686 (224, 224)		20	849.987	0.4803128	0.3507696	0.7615864	0.96050467
0.00275423	13.6704 (224, 224)		20	865.269	0.3602346	0.2630772	0.5711898	0.97040678

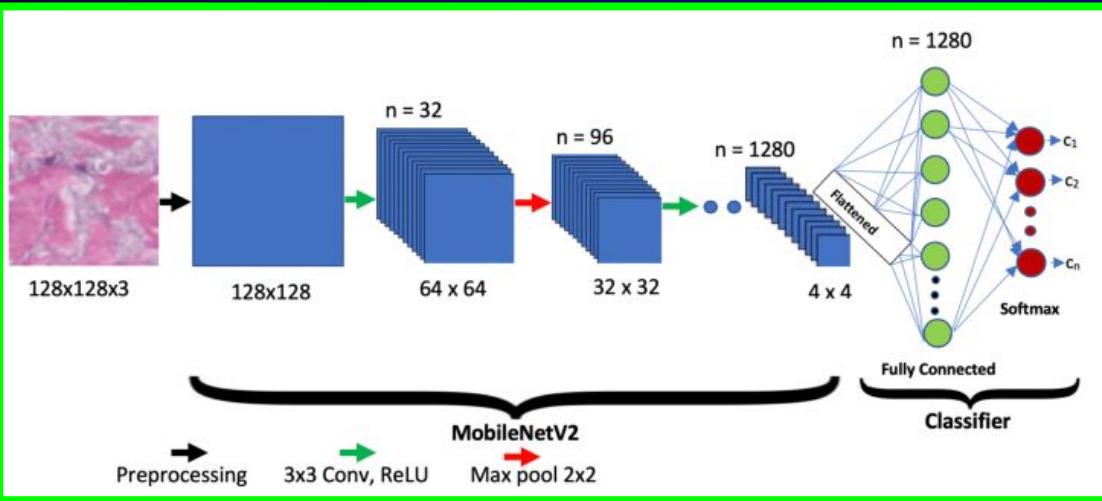
Results using raw CTA as Input

Training Results

lr	pos_weight	resize_dims	iter	total_time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
0.000855178	16.6069 (224, 224)		1	79.4564	0.0598473	0.154839	0.0370943	0.987841
0.00132556	13.5432 (224, 224)		20	862.509	0.600391	0.438462	0.951983	0.990211
0.00199256	12.2775 (224, 224)		20	859.569	0.57637536	0.42092352	0.91390368	0.95060256
0.00109437	16.3686 (224, 224)		20	849.987	0.4803128	0.3507696	0.7615864	0.96050467
0.00275423	13.6704 (224, 224)		20	865.269	0.3602346	0.2630772	0.5711898	0.97040678

Results using raw CTA as Input

Model 3: MobileNetV2



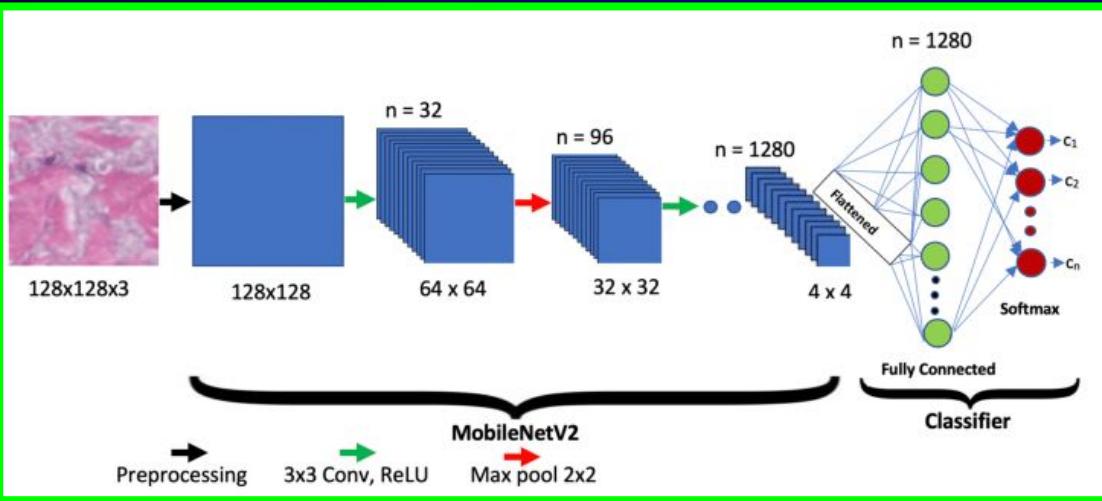
Model Architecture

Tuned Hyperparameters

- Transfer learning: use pre-trained MobileNetV2 architecture trained on Imagenet as feature extractor
- Add classification layer on top
- Make all layers trainable; finetuning MobileNetV2 architecture

```
1 hyperparameters:  
2     input_image_size: (224, 224)  
3     batch_size: 32  
4     lr: 0.0015634  
5     pos_weight: 15.9876  
6
```

Model 3: MobileNetV2



```
1 hyperparameters:  
2     input_image_size: (224, 224)  
3     batch_size: 32  
4     lr: 0.0015634  
5     pos_weight: 15.9876  
6
```

Model Architecture

Tuned Hyperparameters

- Transfer learning: use pre-trained MobileNetV2 architecture trained on Imagenet as feature extractor
- Add classification layer on top
- Make all layers trainable; finetuning MobileNetV2 architecture

Training Results

lr	pos_weight	resize_dims	iter	total_time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
0.000674431	12.7152 (224, 224)	20		589.63	0.770024	0.650438	0.943503	0.993566
0.000700611	15.5943 (224, 224)	20		580.24	0.789864	0.678954	0.955524	0.994589
0.00101701	14.1608 (224, 224)	20		583.94	0.7315228	0.6179161	0.89632785	0.9438877
0.00327364	16.4515 (224, 224)	20		586.57	0.6930216	0.5853942	0.8491527	0.9438877
0.00332455	17.0229 (224, 224)	20		596.54	0.5390168	0.4553066	0.6604521	0.8942094

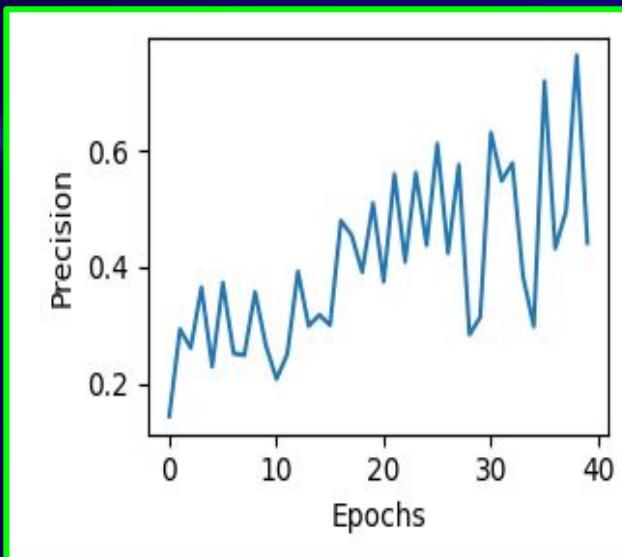
Results using raw CTA as Input

Training Results

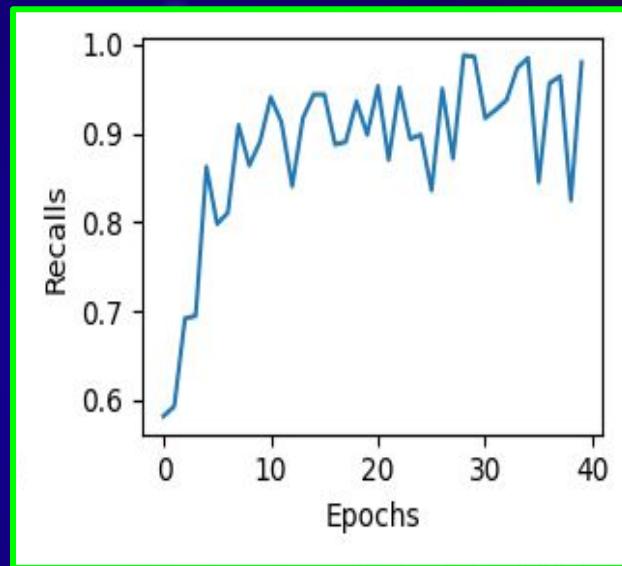
lr	pos_weight	resize_dims	iter	total_time (s)	Val_F1	Val_Precision	Val_Recall	Val_Accuracy
0.000674431	12.7152 (224, 224)	20		589.63	0.770024	0.650438	0.943503	0.993566
0.000700611	15.5943 (224, 224)	20		580.24	0.789864	0.678954	0.955524	0.994589
0.00101701	14.1608 (224, 224)	20		583.94	0.7315228	0.6179161	0.89632785	0.9438877
0.00327364	16.4515 (224, 224)	20		586.57	0.6930216	0.5853942	0.8491527	0.9438877
0.00332455	17.0229 (224, 224)	20		596.54	0.5390168	0.4553066	0.6604521	0.8942094

Results using raw CTA as Input

Analysis of Results: Precision & Recall

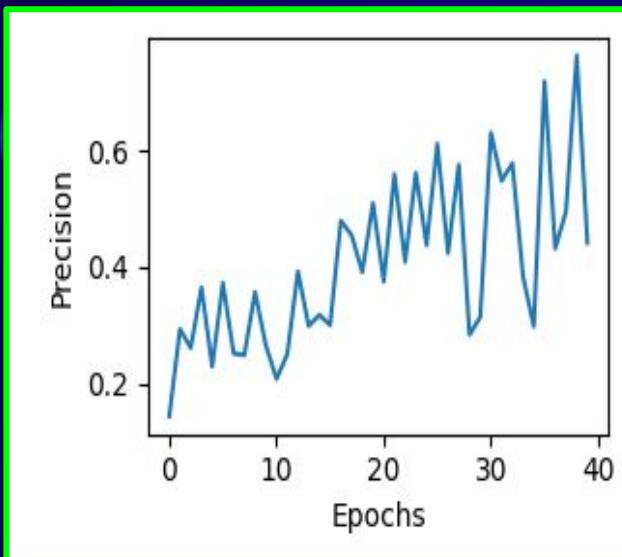


Precision

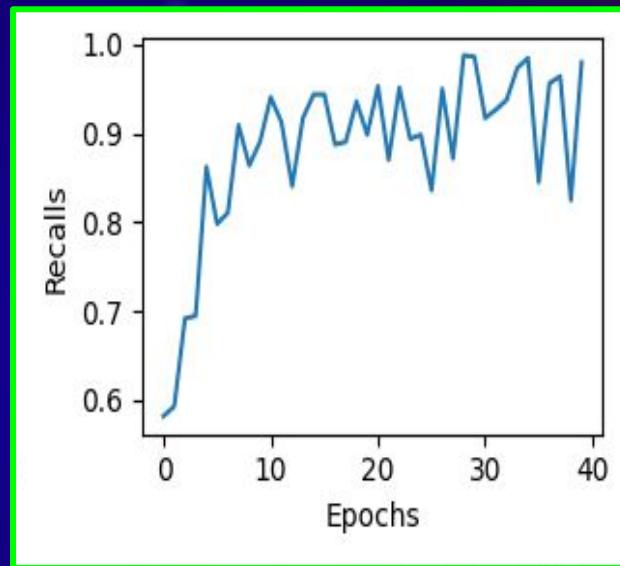


Recall

Analysis of Results: Precision & Recall

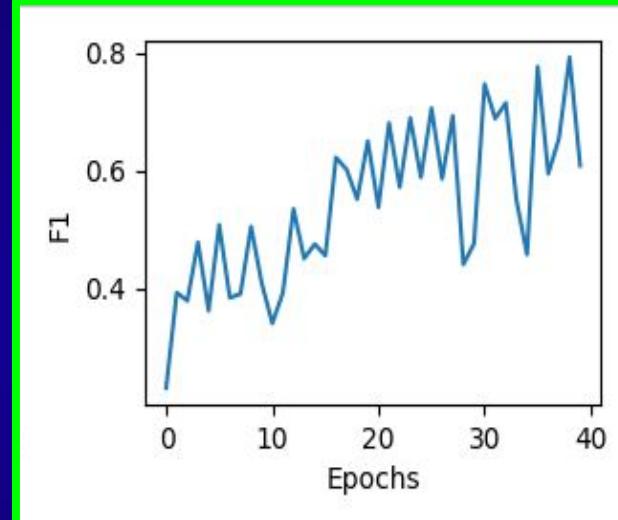
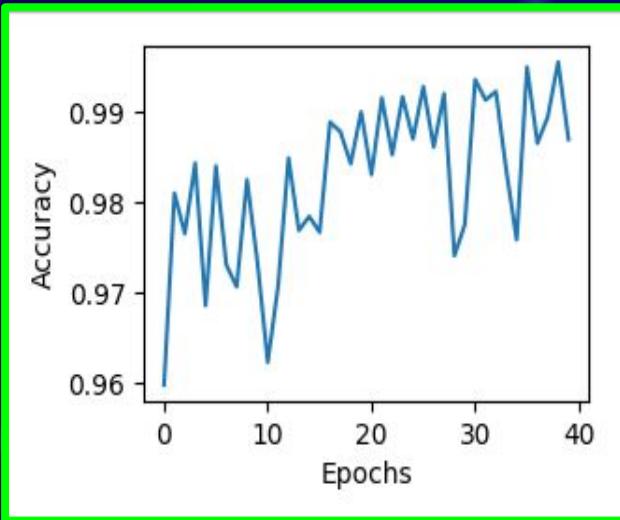


Precision



Recall

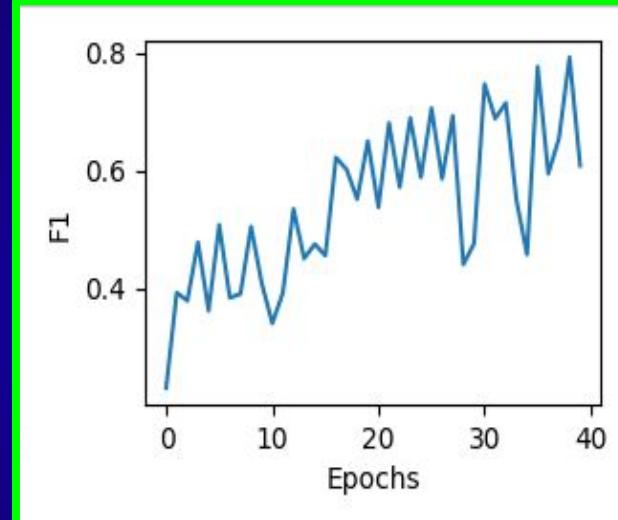
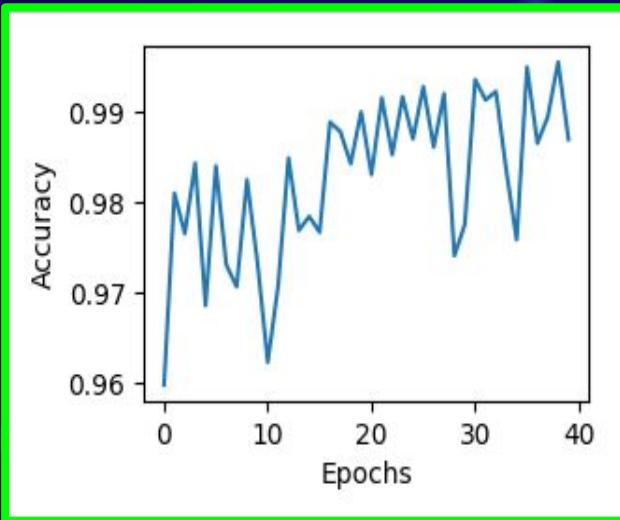
Analysis of Results: Accuracy & F1 Score



Accuracy

F1 Score

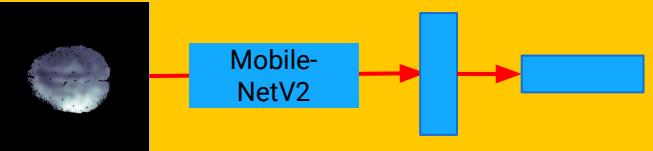
Analysis of Results: Accuracy & F1 Score



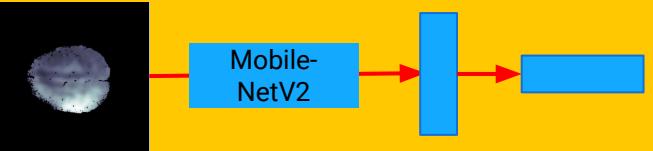
Accuracy

F1 Score

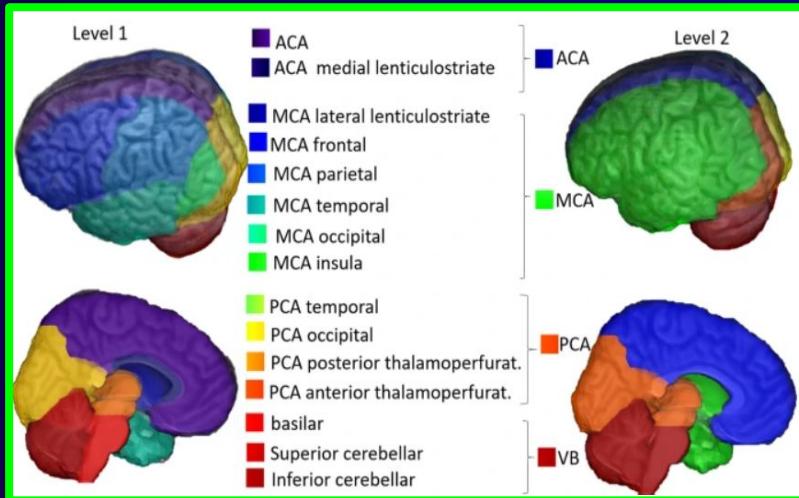
Summary of Results

Approach	Diagram	Best Result
Thresholding SPIRAL and extracting corresponding brain regions		F1: 33.69%, Prec: 32.8%, Rec: 35.29%
Train a Simple CNN		F1: 66.64%, Prec: 51.41%, Rec: 94.69%
Fine-tuning MobileNetV2 as feature extractor		F1: 78.99%, Prec: 67.90%, Rec: 95.55%
Fine-tuning EfficientFormerV2 as feature extractor		F1: 60.03%, Prec: 43.85%, Rec: 95.20%

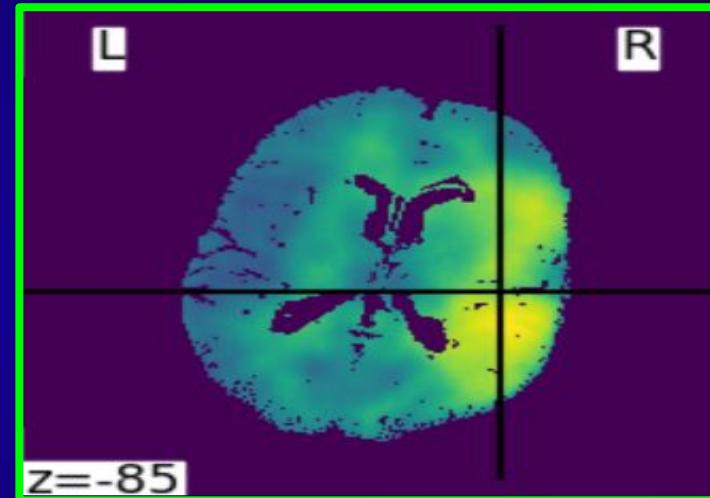
Summary of Results

Approach	Diagram	Best Result
Thresholding SPIRAL and extracting corresponding brain regions		F1: 33.69%, Prec: 32.8%, Rec: 35.29%
Train a Simple CNN		F1: 66.64%, Prec: 51.41%, Rec: 94.69%
Fine-tuning MobileNetV2 as feature extractor		F1: 78.99%, Prec: 67.90%, Rec: 95.55%
Fine-tuning EfficientFormerV2 as feature extractor		F1: 60.03%, Prec: 43.85%, Rec: 95.20%

Discussion: Usefulness

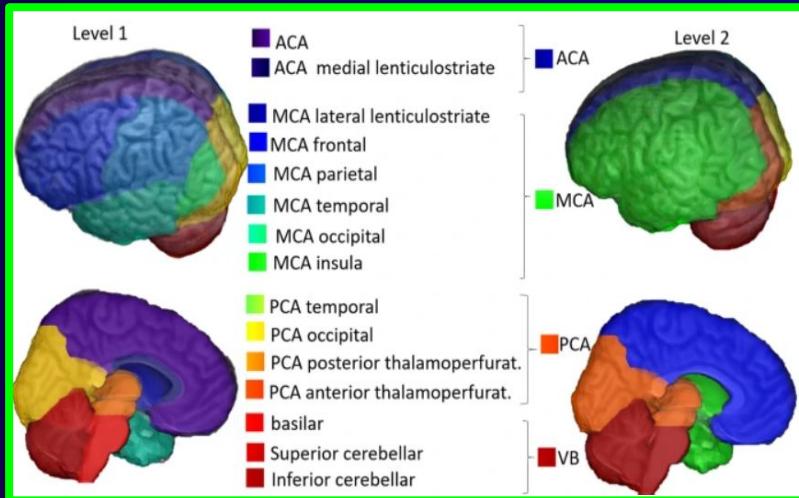


Output = regions

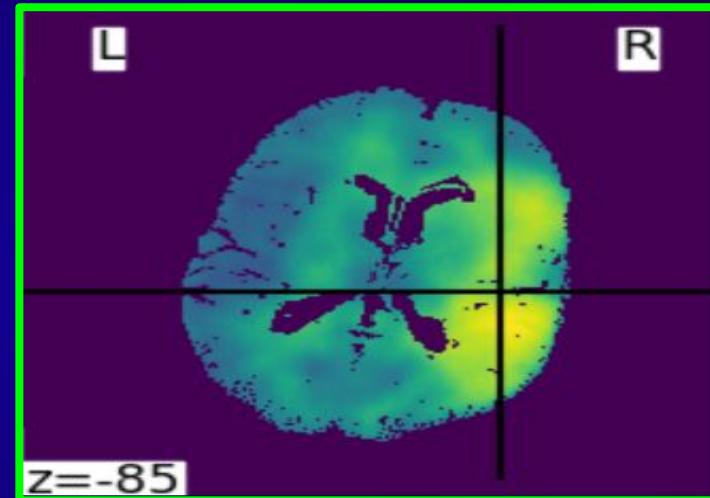


Spiral Scans (AMI)

Discussion: Usefulness

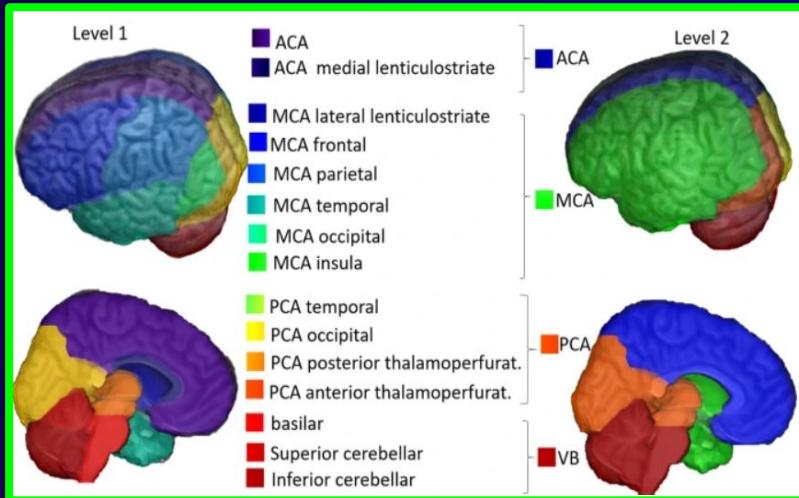


Output = regions

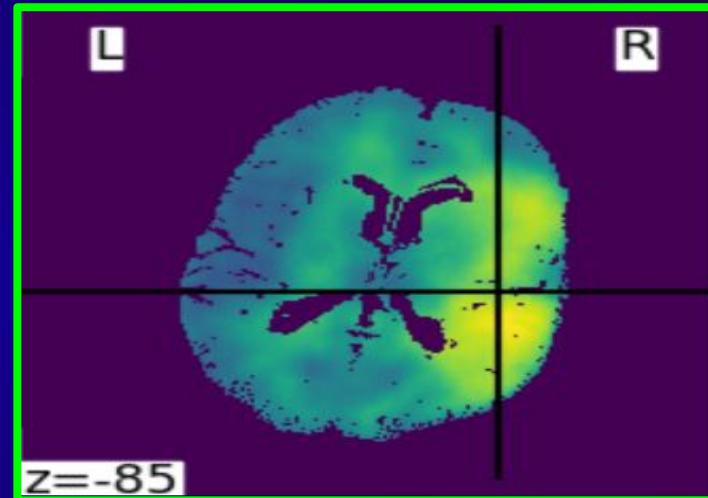


Spiral Scans (AMI)

Discussion: Limitations

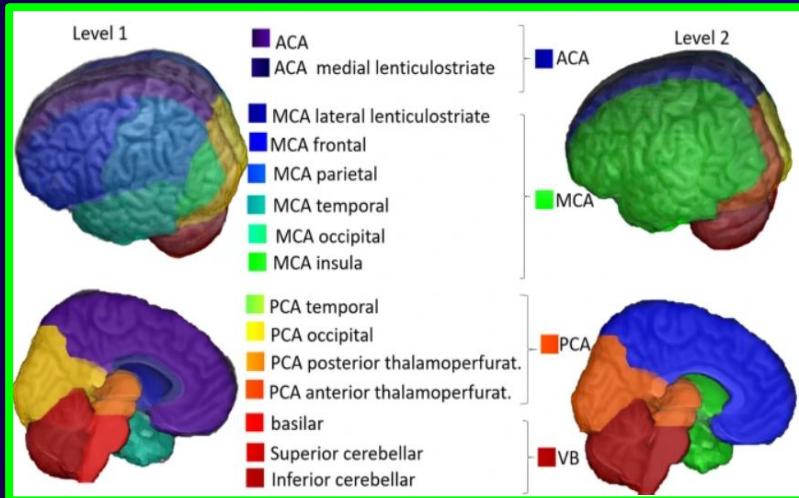


Output = regions

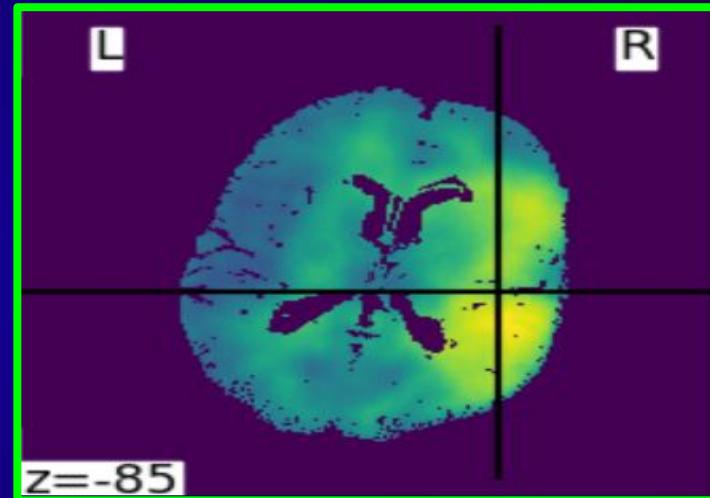


Spiral Scans (AMI)

Discussion: Limitations



Output = regions



Spiral Scans (AMI)

Discussion: Industry Competitors



Discussion: Industry Competitors



**THANKS
FOR
WATCHING**

